



Chapter 5

**Large and Fast:
Exploiting Memory
Hierarchy**

Principle of Locality

→ ram এর উপরে
 cache
 ram } 1 level
 • multiple level
 cache ইত্যাদি

- Programs access a small proportion of their address space at any time
- Temporal locality principle-1
 - Items accessed recently are likely to be accessed again soon
 - e.g., instructions in a loop, induction variables
- Spatial locality principle-2
 - Items near those accessed recently are likely to be accessed soon
 - E.g., sequential instruction access, array data

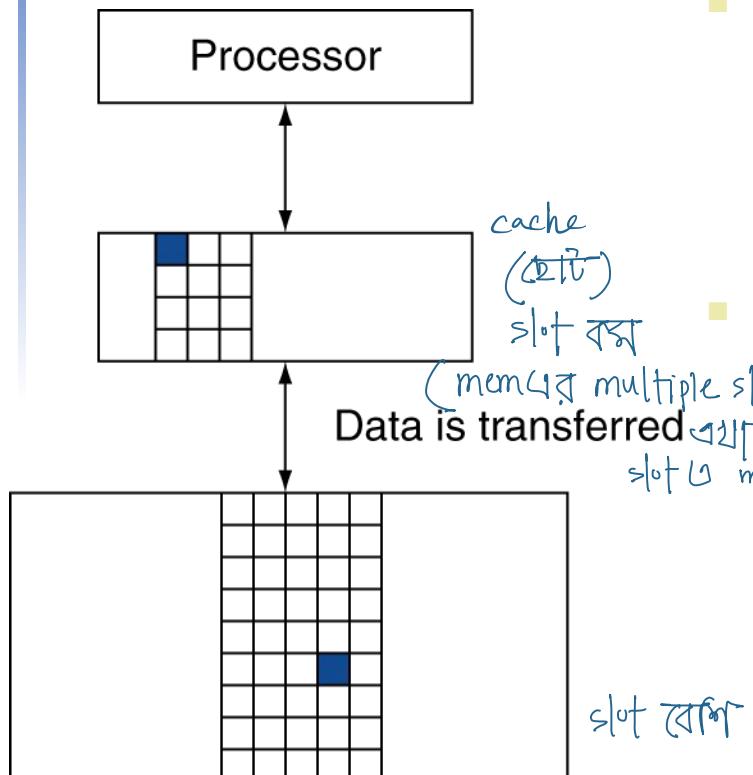


Taking Advantage of Locality

- Memory hierarchy
- Store everything on disk
- Copy recently accessed (and nearby) items from disk to smaller DRAM memory
 - Main memory
- Copy more recently accessed (and nearby) items from DRAM to smaller SRAM memory
 - Cache memory attached to CPU

Memory Hierarchy Levels

- Block (aka line): unit of copying
 - May be multiple words
- If accessed data is present in upper level
 - Hit: access satisfied by upper level
 - Hit ratio: hits/accesses
- If accessed data is absent
 - Miss: block copied from lower level
 - Time taken: miss penalty
 - Miss ratio: misses/accesses
 $= 1 - \text{hit ratio}$
 - Then accessed data supplied from upper level

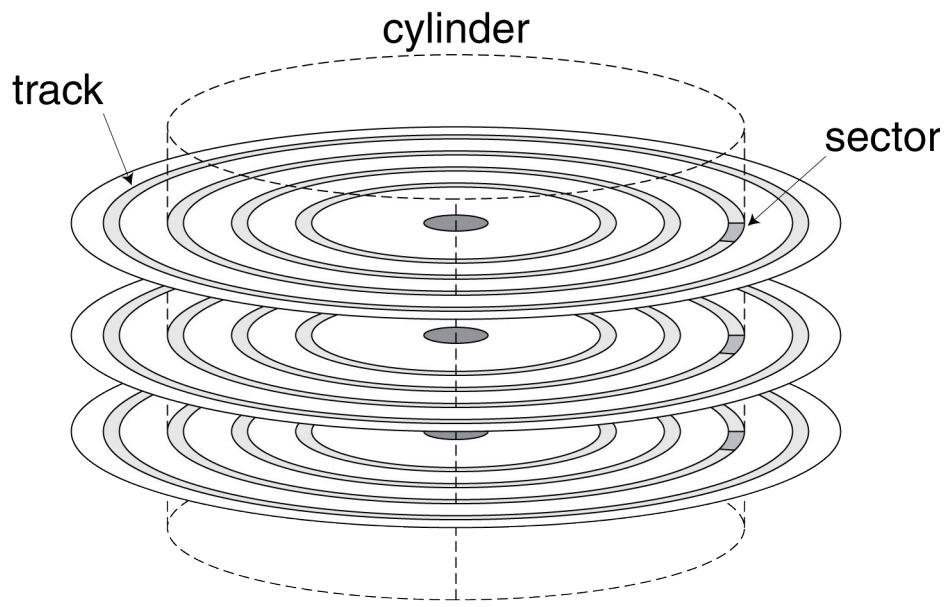


Memory Technology

- Static RAM (SRAM)
 - 0.5ns – 2.5ns, \$2000 – \$5000 per GB
- Dynamic RAM (DRAM)
 - 50ns – 70ns, \$20 – \$75 per GB
- Magnetic disk
 - 5ms – 20ms, \$0.20 – \$2 per GB
- Ideal memory
 - Access time of SRAM
 - Capacity and cost/GB of disk

Disk Storage

- Nonvolatile, rotating magnetic storage



Disk Sectors and Access

- Each sector records
 - Sector ID
 - Data (512 bytes, 4096 bytes proposed)
 - Error correcting code (ECC)
 - Used to hide defects and recording errors
 - Synchronization fields and gaps
- Access to a sector involves
 - Queuing delay if other accesses are pending
 - Seek: move the heads
 - Rotational latency
 - Data transfer
 - Controller overhead

Disk Access Example

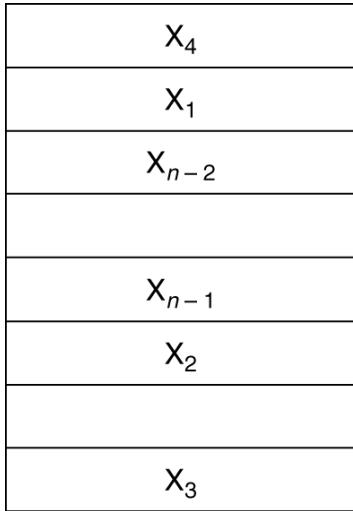
- Given
 - 512B sector, 15,000rpm, 4ms average seek time, 100MB/s transfer rate, 0.2ms controller overhead, idle disk
- Average read time
 - 4ms seek time
 - + $\frac{1}{2} / (15,000/60) = 2\text{ms}$ rotational latency
 - + $512 / 100\text{MB/s} = 0.005\text{ms}$ transfer time
 - + 0.2ms controller delay
 - = 6.2ms
- If actual average seek time is 1ms
 - Average read time = 3.2ms

• seektime minimum ৱ্য hashing ৱ্য

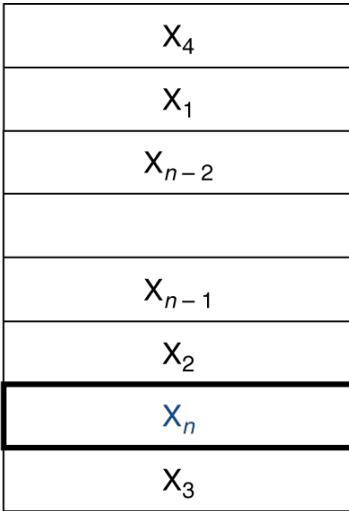


Cache Memory

- Cache memory
 - The level of the memory hierarchy closest to the CPU
- Given accesses X_1, \dots, X_{n-1}, X_n



a. Before the reference to X_n

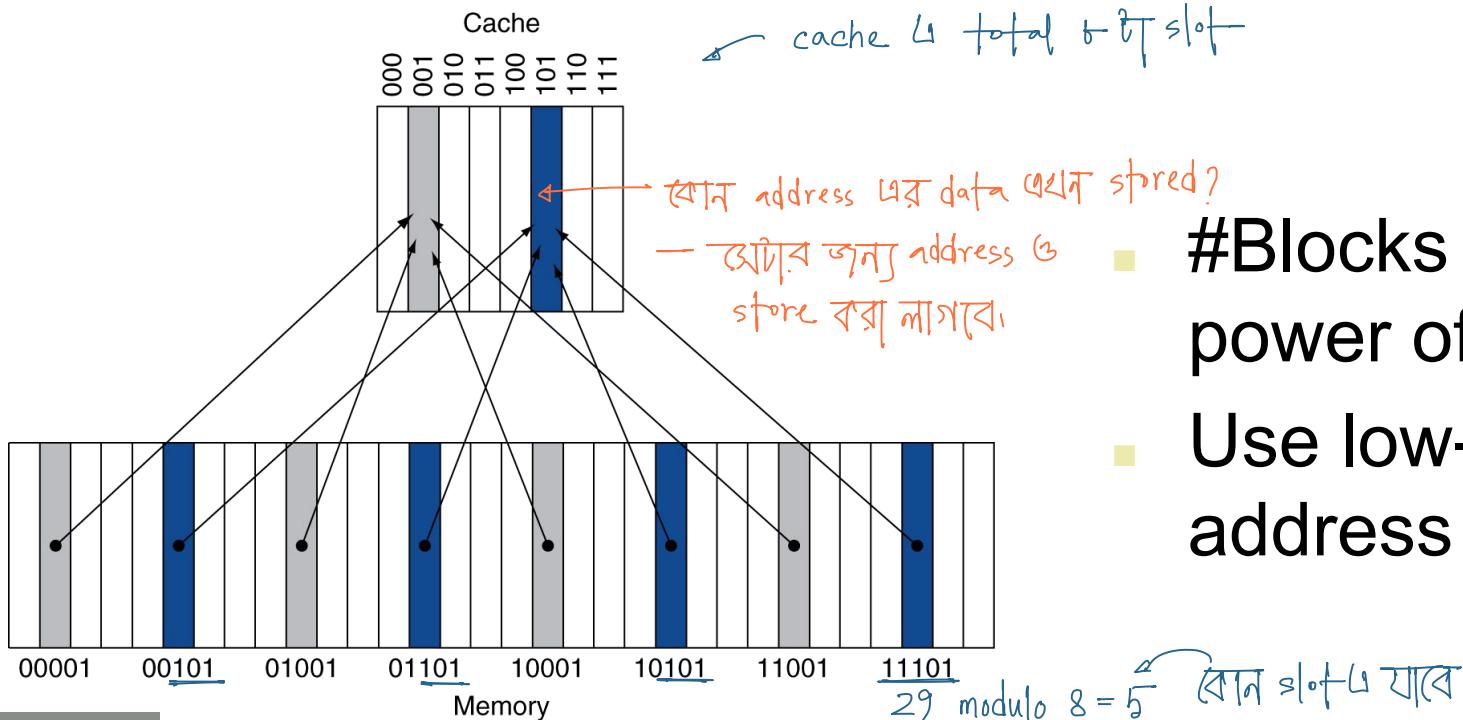


b. After the reference to X_n

- How do we know if the data is present?
- Where do we look?

Direct Mapped Cache

- Location determined by address
- Direct mapped: only one choice
 - (Block address) modulo (#Blocks in cache)



- #Blocks is a power of 2
- Use low-order address bits

minimum time ৰাবণা মডুলো দ্বাৰা কোনো পথে যেন এটিৱে অন্তৰ জন্মে clock cycle না থাকিব।

- ALU use কৰিবো

$$x \% 2^n \rightarrow x \& (2^n - 1)$$

$$\begin{array}{r} 0 \ 1 \ 1 \ 1 \dots \\ \hline n \text{ মসৃষ্ট } 1 \end{array}$$

সুতৰে last n bits পাইবা

$2^n \rightarrow$ cache কৰিবলৈ block আকৰণ

last n bit দিয়ে \rightarrow index কৰিবলাব

\hookrightarrow যেমন address map হয় তদেৱে ইই n bits same কৰাৰ।

সুতৰে store কৰিবলৈ higher bit মূলো — tag bit কৰলো (uniquely identify indexing কৰিবলৈ কৰিবলৈ upper bit) \Rightarrow কৰাবলৈ আকৰণ

valid — actually data এখন store কৰা আছে বিনা / valid বিনা।

Tags and Valid Bits

- How do we know which particular block is stored in a cache location?
 - Store block address as well as the data
 - Actually, only need the high-order bits
 - Called the tag
- What if there is no data in a location?
 - Valid bit: 1 = present, 0 = not present
 - Initially 0

Cache Example

- 8-blocks, 1 word/block, direct mapped
- Initial state

address [13]

22, decimal
value

26,
22,
26,
16,
3,
16,
18,
16

	Index	V	Tag	Data
000	N			
001	N			
010	N			
011	N			
100	N			
101	N			
110	N			
111	N			



Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Miss	110

22,

26,

22,

26,

16,

3,

16,

18,

16

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

valid - N

so Miss

memory থেকে read করে

আনাগুর
then valid
করে দিয়ে

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
26	11 010	Miss	010

22,

26,

22,

26,

16,

3,

16,

18,

16

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Hit	110
26	11 010	Hit	010

22,

26,

22,

26,

16,

3,

16,

18,

16

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

valid - yes

Tag - same

Hit

— আগেই load

করা আছে।

Cache Example

22,

26,

22,

26,

16,

3,

16,

18,

16

Word addr	Binary addr	Hit/miss	Cache block
16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Cache Example

Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

22,

26,

22,

26,

16,

3,

16,

18,

16



Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Cache Example

Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

22,

26,

22,

26,

16,

3,

16,

18,

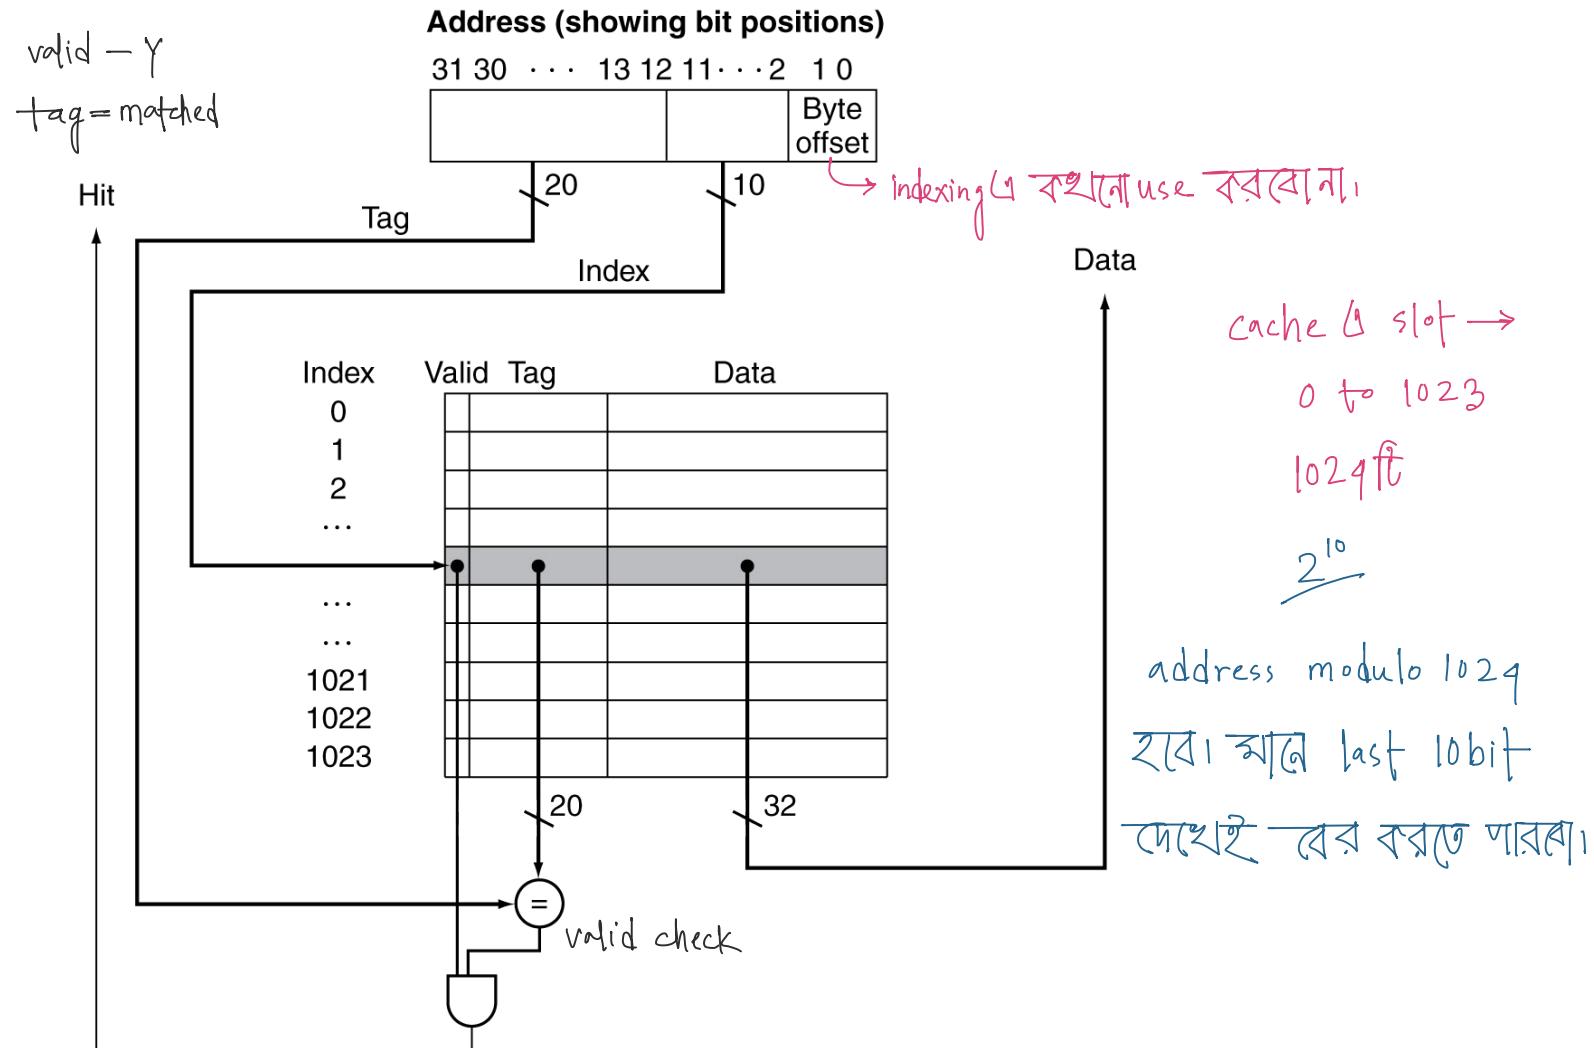
16



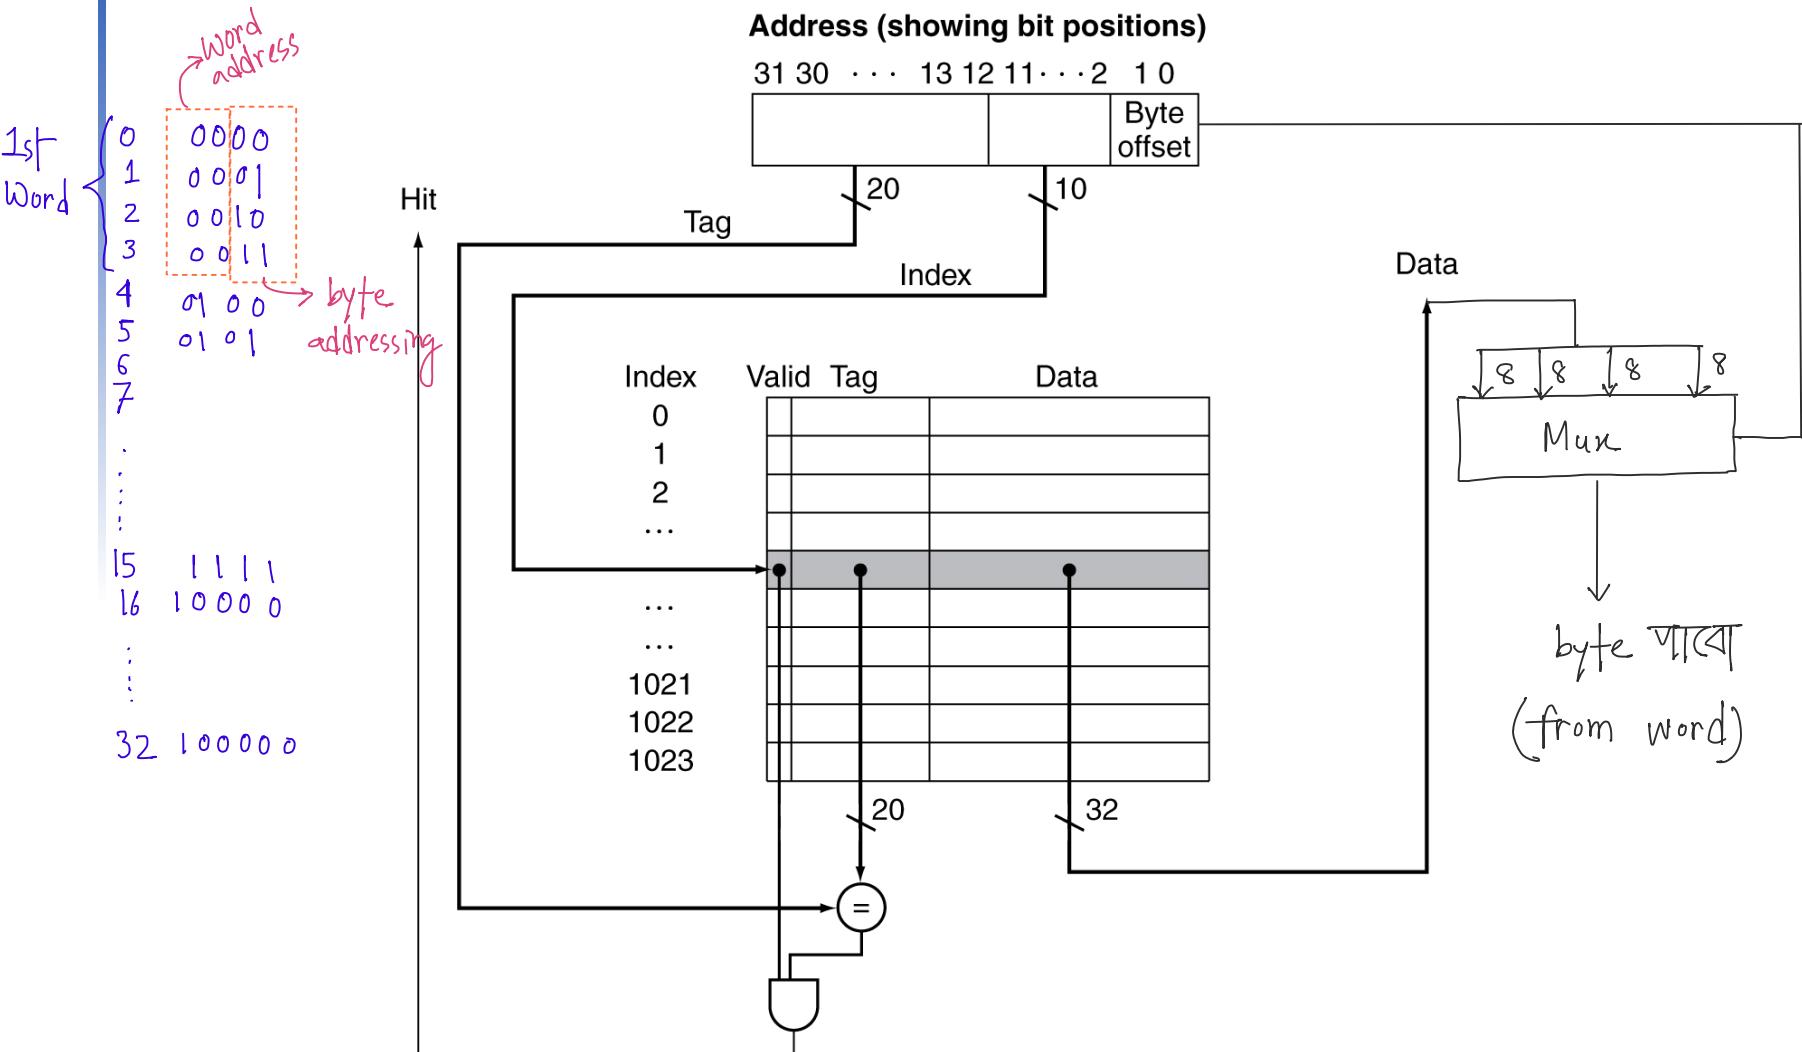
Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	10	Mem [10010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Address Subdivision

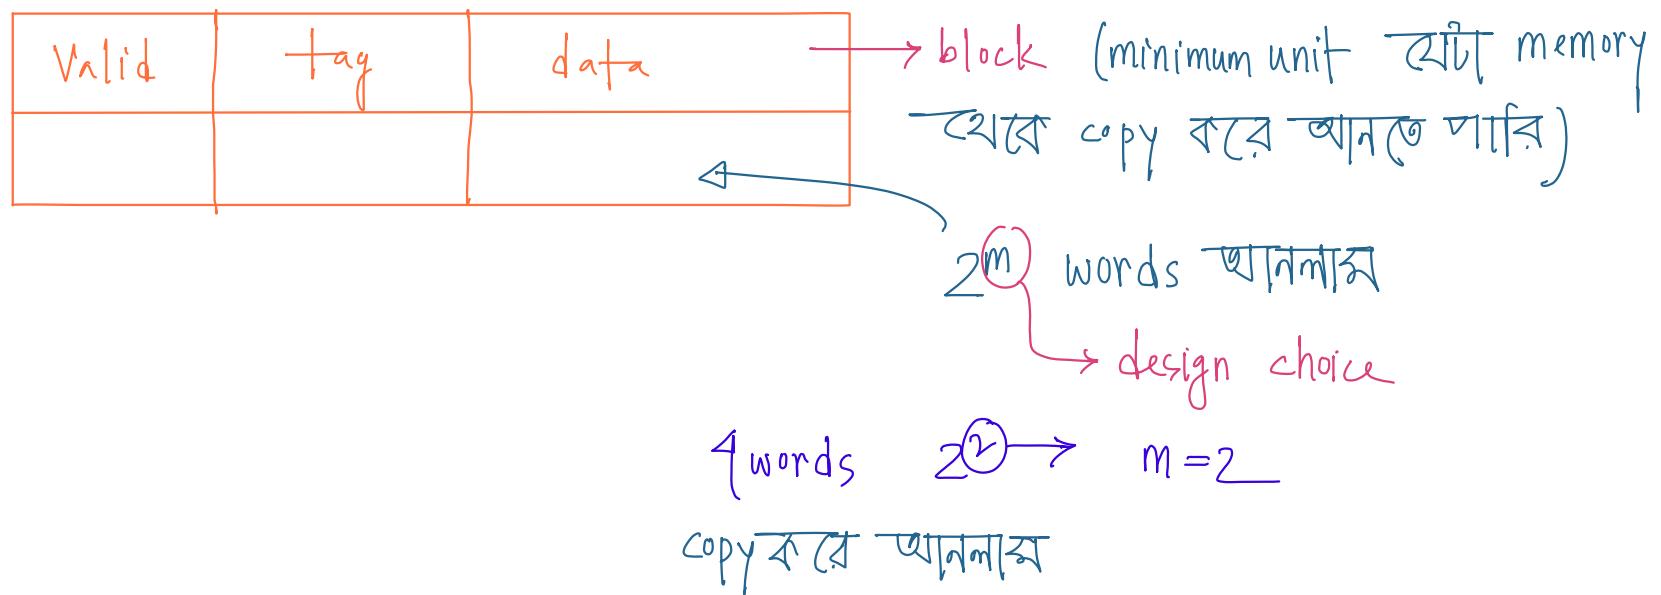


Address Subdivision



byte
 read → 0 0 0 0 0 }
 0 0 0 0 1 } 1 word
 0 0 0 1 0 } next word address
 0 0 0 1 1 }
 0 0 1 0 0

read → 0 0 1 0 0



10 bit এর last 2 bit ignore → byte এর জন্য

7th - 8th bit → word index

0 - 5 bit যদি নেই তবে 4 bit word এ uniquely identify করা থাবে

10 10 1 00 0
 0 1
 1 0
 1 1

1st word

10 10 1 1 0 0 2nd word

1011 0 0 00 → 4 word পারে চল আমলাব্দ

8 words copy করে থানাতে 3 bit লাগতে।

Tag size:
 $32 - (n + m + 2)$
 indexing
 word
 byte offset

Total bits in a cache

Size of tag field

v is the valid field size, i.e., 1

- 32-bit addresses
- A direct-mapped cache
- The cache size is 2^n blocks
 - so n bits are used for the index
- The block size (b) is 2^m words (2^{m+2} bytes = 2^{m+2+3} bits)
 - m bits are used for the word within the block
 - two bits are used for the byte part of the address

$$2^m \times 4 \text{ bytes} \quad 2^{m+2} \times 8 \text{ bits}$$

- Total number of bits (C)

- $C = 2^n \times (b + t + v)$
- $t = 32 - (n + m + 2)$
- $b = 2^{m+5}$

C

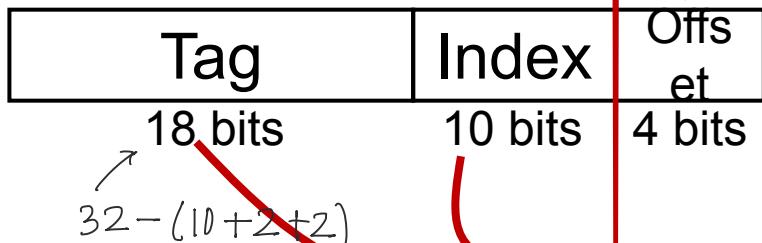
$$\begin{aligned} &= 2^n \times (2^{m+5} + 32 - (n \\ &+ m + 2) + 1) \\ &= 2^n \times (2^m \times 32 + 32 \\ &- n - m - 1) \end{aligned}$$

Total bits in a cache

CT

- How many total bits are required for a direct-mapped cache with 16 KiB of data and 4-word blocks (i.e., 16 bytes), assuming a 32-bit address?

31



data portion size

$$\begin{aligned}16 \text{ KiB} &= 4 \text{ Ki Words} = \\1 \text{ Ki Blks} &= 2^{10} \text{ Blks} \\C &= 1024 \times (b + t + v) \\&= 1024 \times (4 \times 32 + t + 1) \\&= 1024 \times (4 \times 32 + 18 + 1) \\&= 147 \text{ Kilo bits}\end{aligned}$$

16 KiB data $\rightarrow 16 \times 2^{10}$ byte
 $\rightarrow \frac{16 \times 2^{10}}{4}$ words
 $= 4096$ words cache store রয়ে
যাব।

1024 slots

2^{10}

$n = 10$

4 words
 $= 2^2$
 $m = 2$

$$2^{10} * (1 + 18 + 4 \times 32)$$

\uparrow \uparrow \downarrow
valid tag 4 words
 32 bits each.

147 kilobits \rightarrow 18 kilobytes

Tag, Valid bit — overhead

Example: Larger Block Size

- 64 blocks, 16 bytes/block
total cache slots
4 words

- To what block number does address 1200 map?

- Block address = $\lfloor 1200/16 \rfloor = 75$

- Block number = 75 modulo $\underline{64} = 11$

31 10 9 4 3 0

Tag	Index	Offset
-----	-------	--------

22 bits

6 bits

64 blocks

4 bits

16 bytes

(last 2 bits) byte address = 2 bits

word address = 2 bits

Byte Address
প্রথম block
address এ যাওয়া
লাগবে।

করণ
1200-1215 no
byte same
block টি
যাববে।
floor ফিটি
block by block
copy। সতি
block (16).

1210 \rightarrow 11 টি আববে।

11 টি 16 bytes আববে

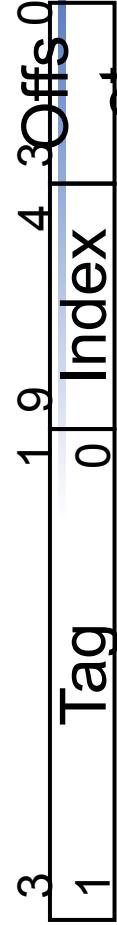
1210 আববে 10 no. bytes।

Copy করণ সময় 16 bytes বয়টি।

block পুরোটি পড়া লাগবে। 1210

(10) টেক্স mux use করা লাগবে এবং

Example: Larger Block Size



- 64 blocks, 16 bytes/block
 - To what block number does address 1200 map?

$$\text{Block address} = \lfloor \frac{1200}{16} \rfloor = 75$$

$$\text{Block number} = 75 \bmod 64 = 11$$

- In fact Block 11 maps all addresses between 1200 and 1215.

$$\left\lfloor \frac{\text{Byte address}}{\text{Bytes per block}} \right\rfloor \times \text{Bytes per block}$$

$$\left\lfloor \frac{\text{Byte address}}{\text{Bytes per block}} \right\rfloor \times \text{Bytes per block} + (\text{Bytes per block} - 1)$$

Byte Address
↑

[byte address /
bytes per block]

Block Size Considerations

- Larger blocks should reduce miss rate
 - Due to spatial locality — যেনি word আনতে পারবো।
- But in a fixed-sized cache *let 4 থেকে 16 ব্যৱহাৰ
— 12টা word যেনি বাধ্যতা
পারবো। miss rate বৃদ্ধয়ে।*
 - Larger blocks \Rightarrow fewer of them
 - More competition \Rightarrow increased miss rate
- Larger miss penalty
 - Larger blocks \Rightarrow Larger transfer time
 - Can override benefit of reduced miss rate
 - Early restart critical-word-first can help



Block Size Considerations

- Larger blocks should reduce miss rate
 - Due to spatial locality
- But in a fixed-sized cache
 - Larger blocks \Rightarrow fewer of them
 - More competition \Rightarrow increased miss rate
- Larger miss penalty
 - Larger blocks \Rightarrow Larger transfer time
 - Can override benefit of reduced miss rate
 - Early restart critical-word-first can help

→ cache size fixed
block size বাড়ালায়
এক ম্যাপের ক্ষেত্রে
multiple address
same slot এ map.
→ miss rate ঘোড়ে গলা।

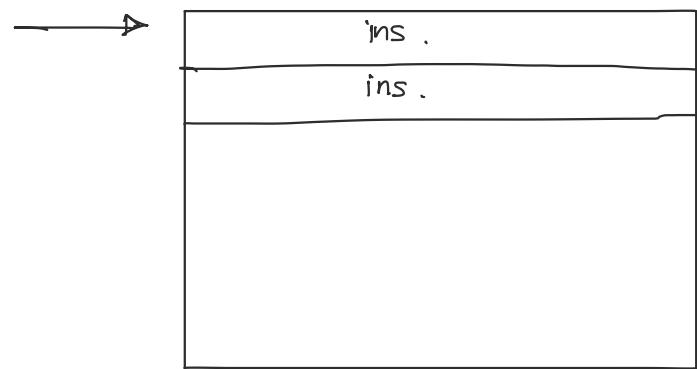
→ এশি data
memory থেকে
copy করে আনতে
মজবুত এশি
গোপনীয়।
→
net benefit করে গলা।



Early restart

→ এ word টি দ্রবার ঘোটা চলে আগলে CPU execution resume করবো। Parallelly যাবিটা loading হবে।
→ pipeline এ still লাগে মাধ্যম memory থেকে data আসা আছে।

- resume execution as soon as the requested word of the block is returned; Does not wait for the entire block
- For instruction access, it works best
 - Instruction accesses are largely sequential
 - so if the memory system can deliver a word every clock cycle, the processor may be able to restart operation when the requested word is returned, with the memory system delivering new instruction words just in time.
 - This technique is usually less effective for data caches because it is likely that the words will be requested from the block in a less predictable



ins. execute \leq one clock cycle M[ST]
best case: sequential instruction

\downarrow \rightarrow block copy $\overline{22}$

Critical Word First

→ যে word দাঁ গাই যেটা আগে copy করে
আপনার প্রসেসরে CPU resume হবে।

- Organizes the memory so that the requested word is transferred from the memory to the cache first.
- The remainder of the block is then transferred, starting with the address after the requested word and wrapping around to the beginning of the block.
- Can be slightly faster than early restart
- but it is limited by the same properties that limit early restart.

Same problem: data access এবং predictable
নাহীন।



Cache Misses

- On cache hit, CPU proceeds normally
- On cache miss
 - Stall the CPU pipeline
 - Fetch block from next level of hierarchy
 - Instruction cache miss
 - Restart instruction fetch
 - Data cache miss
 - Complete data access



Write-Through

→ cache ৱি hit হলে cache and memory SW update করবে

- On data-write hit, could just update the block in cache
 - But then cache and memory would be inconsistent

Write through: also update memory

- But makes writes take longer
 - e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
 - Effective CPI = $1 + 0.1 \times 100 = \underline{11}$

Solution: write buffer

→ খুব বেশি write পায়লে buffer full হবে।
তখনকি same problem. Buffer full
হলে memory তে write so stall করা
পারবে।

- Holds data waiting to be written to memory
- CPU continues immediately
 - Only stalls on write if write buffer is already full

Write-Back

— cache গু update, memory টি করবো না,
 → lazy reflection হবে।
— dirty bit flag থাকবে।

- Alternative: On data-write hit, just update the block in cache
 → dirty bit ১ করবো

- Keep track of whether each block is dirty

- When a dirty block is replaced

- Write it back to memory
 - Can use a write buffer to allow replacing block to be **read first**
 → হেটা নতুন ঘৰচে
 মেটাই priority
 → পাৰে।

memory reflection when?

→ যখন cache গু মেটই address
 replace হবে।

	x20	1
--	-----	---

dirty bit → 1
 ↑
cache গু updated
memory টি ৱি নি।

write এ miss \rightarrow Write through scheme \rightarrow memory, cache (update)

allocate on miss \rightarrow memory টি update

-যেটাই আবার cache এ যানযো।

cause -যেটা আবার ই এরতে পারিব।

Cache এ load করতে পারিব। \rightarrow e.g. initialization এর case এ ক্ষেত্রে

শুধু sw

sw

-থার্ড

Write Allocation

- What should happen on a write miss?
- Alternatives for write-through
 - Allocate on miss: fetch the block ↗ Write allocate
 - Write around: don't fetch the block
- No Write ↗ ■ Since programs often write a whole block before reading it (e.g., initialization)
- For write-back
 - Usually fetch the block
 - (See next slides)



Advantage for write-through

- we can write the data into the cache and then read the tag;
 - tag bit compare করবেই update করলাগ্য হিন্দু
গোটা অ্যামেলি cache miss হিল।
- if the tag mismatches, then a miss occurs.
- Because the cache is write-through, the overwriting of the block in the cache is not catastrophic
 - memory has the correct value and we can do it right.
- THIS CANNOT BE DONE FOR WRITE BACK → write back & lazily update হয়। এই প্রিং বাজে করবে না। dirty bit 1 হিল...



Write Back

- If we have a cache miss, we must first write the block back to memory if the data in the cache is modified.

✓ stores require two cycles:

- a cycle to check for a hit
- followed by a cycle to actually perform the write

- Alternative: Write Buffer to hold that data

— যদি cache এ write করবে তবুও buffer এ রেখে CPU execution

resume করবো, hit হলে পরে এখন থেকে write করবে।

(pipelining করে এক cycle এ করতে চাইলে) ↗

(next clock cycle এ)
unused



Write Back- Write Buffer

- write buffer holds the data to write in effectively allowing the store to take only one cycle by pipelining it:
 - When a store buffer is used, the processor does the cache lookup and places the data in the store buffer during the normal cache access cycle.
 - Assuming a cache hit, the new data is written from the store buffer into the cache on the next unused cache access cycle.

Write Back- Write Buffer for miss

- the modified block is moved to a write-back buffer associated with the cache in case of a miss
 - while the requested block is read from memory.
 - The write-back buffer is later written back to memory.
- Assuming another miss does not occur immediately, this technique reduces the miss penalty

পর্যবেক্ষণ অনেকগুলো miss —স্টল দুওয়া নাগে

buffer পর্যবেক্ষণ করা নিষেধ হবে।



Example: Intrinsity FastMATH

processor

- Embedded MIPS processor
 - 12-stage pipeline
 - Instruction and data access on each cycle
- Split cache: separate I-cache and D-cache
 - Each 16KB: 256 blocks × 16 words/block
 - D-cache: write-through or write-back
- SPEC2000 miss rates
 - I-cache: 0.4%
 - D-cache: 11.4%
 - Weighted average: 3.2%



Example: Intrinsity FastMATH

16×1 mux
→ selection block offset

$\xrightarrow{\text{words/block}}$

Address (showing bit positions)

$$\begin{aligned} \text{tag bit} &= 32 - 8 - 9 - 2 \\ &= 18 \end{aligned}$$

$2^8 \rightarrow$ 8 bit indexing

$256 \text{ blocks} \times 16 \text{ words/block}$
 $\xrightarrow{\text{block offset}} = 4$

Hit

In practice, to eliminate the multiplexor, caches use two RAMS



a smaller RAM for the tags



a large RAM for the data with the block offset supplying the extra address bits

Measuring Cache Performance

Components of CPU time

- CPU execution cycles
 - Includes cache hit time
- Memory stall cycles
 - Mainly from cache misses

metrices → compare বয়ের
জন্য লাগবে।

→ execution বয়েত data দরবার যেটি
cache এ নাই। memory থেকে
আসা লাগবে।

Measuring Cache Performance

- Read-stall cycles = $\frac{\{\text{Reads}\}}{\{\text{Program}\}} \times$
Read miss rate × Read miss penalty

- Write-stall cycles = $\left(\frac{\{\text{Writes}\}}{\{\text{Program}\}} \times \right.$
Write miss rate × Write miss penalty $) +$
Write buffer stalls

→ write through scheme හි buffer introduce බව දෙයින්। Buffer full

බය ගැලී buffer stall ලැබේ।

For Write-through Cache



Notes: Write Buffer Stalls

- Write buffer stalls depend
 - on the proximity of writes
 - and not just the frequency
 - Usually we can ignore write buffer stalls
 - in systems with 4 or more words depth
 - With a memory capable of accepting writes at a rate that significantly exceeds the average write frequency in programs (by a factor of 2)
 - If a system did not meet these criteria, it would not be well designed;
 - instead, the designer should have used either a deeper write buffer or a write-back organization
- বাছাৰা time
ক'ত দুটি আয়তে
So not easy to deduce an equation
গৱেষণা গৱেষণা

Measuring Cache Performance

- With simplifying assumptions:
 - read and write miss penalties are same
 - In most write-through schemes this is the case
 - Write buffer stalls are negligible

Memory stall cycles

$$= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty}$$


For
Data

Memory stall cycles

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty}$$


For
Instruction

All instruction search

করা লাগে। Cache
বেথায় আছে।

instruction এটি করে
পাও লাগে

Cache Performance Example



Given

- miss rate: I-cache = 2% ; D-cache = 4%
- Miss penalty = 100 cycles
- Base CPI (ideal cache) = 2
- Load & stores are 36% of instructions

How much faster is a processor with a perfect cache?

Say total instruction: I

- I-cache miss cycles: $I \times 0.02 \times 100 = 2.00 \times I$
- D-cache: $I \times 0.36 \times 0.04 \times 100 = 1.44 \times I$

per instruction
2

$$\text{Actual CPI} = 2 + 2 + 1.44 = 5.44$$

Ideal CPU is $5.44/2 = 2.72$ times faster

Amdahl's Law

- A rule stating that the performance enhancement possible with a given improvement is limited by the amount that the improved feature is used.
- What happens if the processor is made faster, but the memory system is not?
 - The amount of time spent on memory stalls will take up an increasing fraction of the execution time

Cache Performance Example 2



Given

- miss rate: I-cache = 2% ; D-cache = 4%
- Miss penalty = 100 cycles
- Base CPI (ideal cache) = 1 Base CPI improved
- Load & stores are 36% of instructions

How much faster is a processor with a perfect cache?

Say total instruction: I

- I-cache miss cycles: $I \times 0.02 \times 100 = 2.00 \times I$
- D-cache: $I \times 0.36 \times 0.04 \times 100 = 1.44 \times I$

Actual CPI = $1 + 2 + 1.44 = 4.44$

Ideal CPU is $4.44/1 = 4.44$ times faster

cache দ্বারা এনে
improvement হয়ে
Processor যাতে পারছে।

miss না হলো \rightarrow super fast

miss হলো \rightarrow performance degrade.

Lesson (re)Learned

- So, in Example 2:
 - we have improved the processor architecture
 - Ideal CPI is 1 (previously it was 2)
 - Memory system was NOT improved
 - Now effective CPI has become 4.44
 - So, there is a small improvement; Not what we expected

Previously
it was 5.44

- The amount of execution time spent on memory stalls:

- Increases From $\frac{3.44}{5.44} = 63\%$ to $\frac{3.44}{4.44} = 77\%$

processor idle
থাকে। স্ট্যাম্প।

cache এর ফলে
যদি আছে।

Performance lost increases

total time
এর 63% cache।

total time এর 77%
cache।

Average Access Time

- Hit time is also important for performance
- Average memory access time (AMAT)
 - $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- Example
 - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, I-cache miss rate = 5%
 - $AMAT = 1 + 0.05 \times 20 = 2\text{ns}$
 - 2 cycles per instruction

Performance Summary

- When CPU performance increased
 - Miss penalty becomes more significant — cache improve
করিন্তাই (but
processor improve
ব্যর্থ)
- Decreasing base CPI
 - Greater proportion of time spent on memory stalls
- Increasing clock rate
 - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance



Associative Caches

→ directly mapped তাও রেট পাবে

memory block & cache block

- Fully associative ↗ একটি address cache এর স্বত্ত্বে যেখানে entry টি যেতে পাবে, fixed mapping নেই।

- Allow a given block to go in any cache entry
- Requires all entries to be searched at once
- **Comparator per entry (expensive)** → parallel যব
জায়গায় থুঁজবে

- n -way set associative → directly mapped, fully associated এর বাবে

- Each set contains n entries
- Block number determines which set
 - (Block number) modulo (#Sets in cache)
- Search all entries in a given set at once
 - n comparators (less expensive)

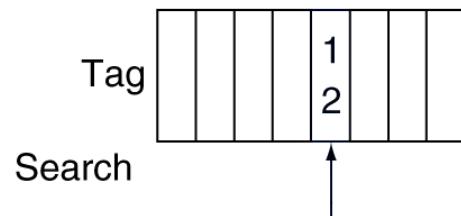
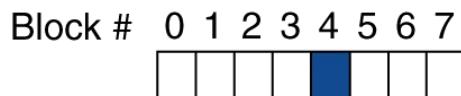
index bit দ্বাৰা
set দ্বীপা determine
কৰবে,

— specific set
search দ্বৰো
cache লাগবে না।

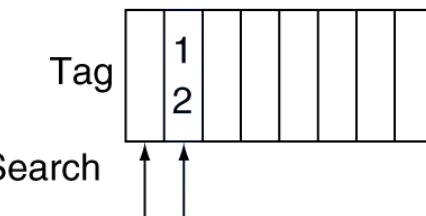
direct map ↗ → cache টি জায়গা হ'কা আকলেও replace কৰা লাগচে। Utilize বৱ্যাছেন।

Associative Cache Example

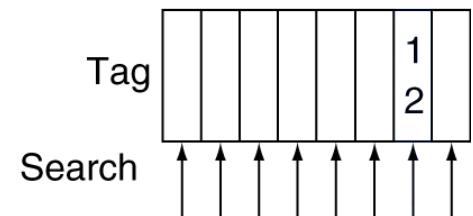
Direct mapped



2 way **Set associative**



Fully associative

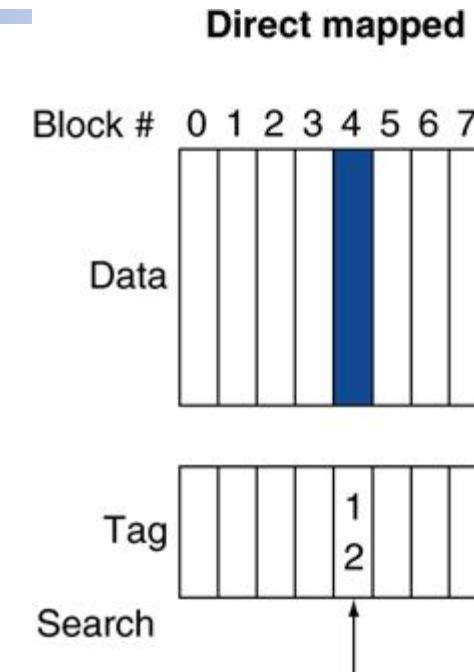
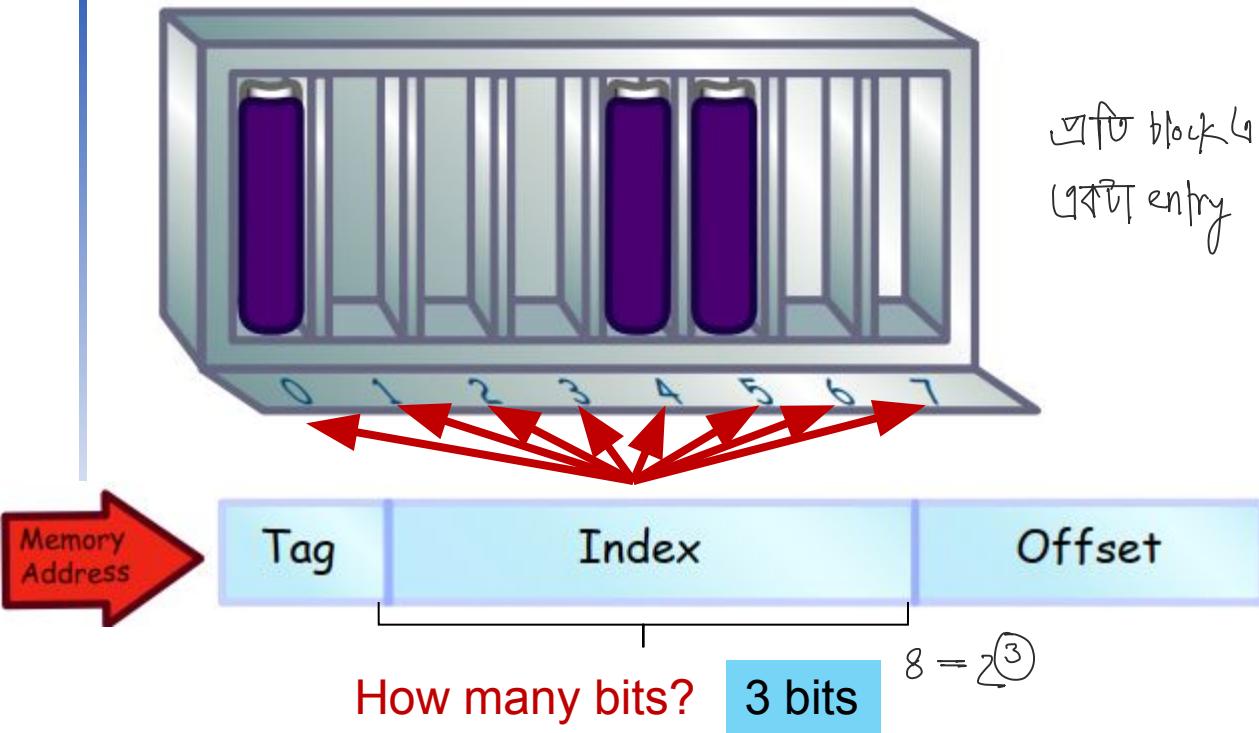


দুর্ধো cache এর মেশিনে

জায়গাময় খবড়ে দারে।

1-Way Set Associative

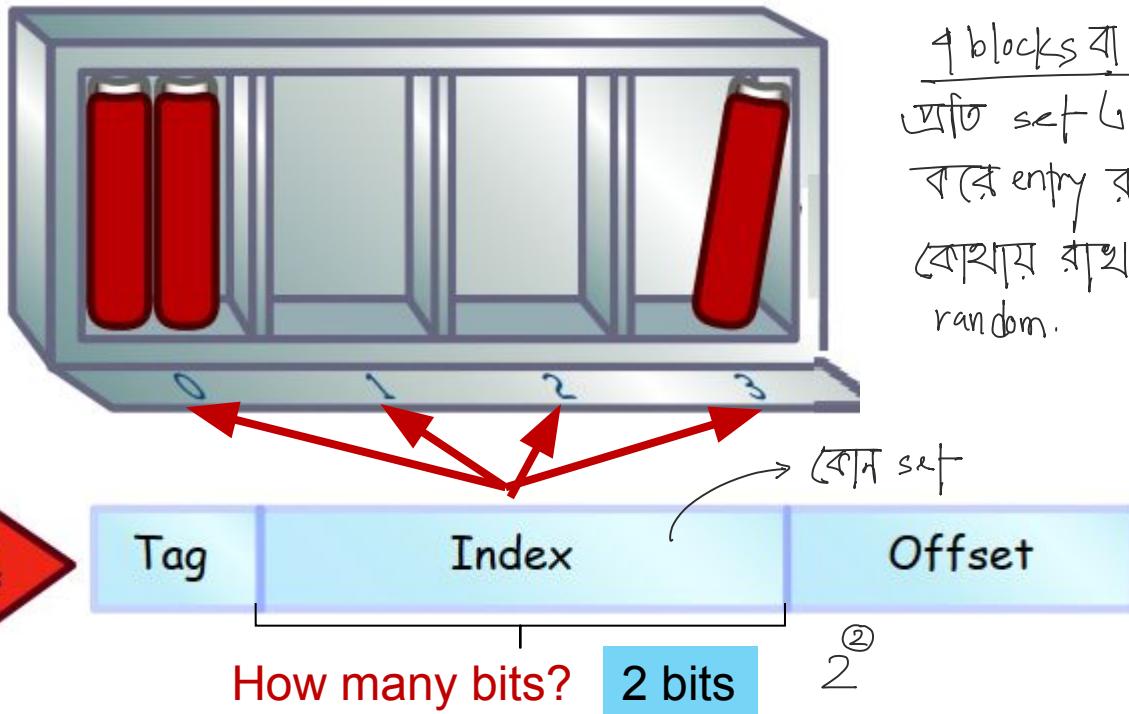
→ direct map



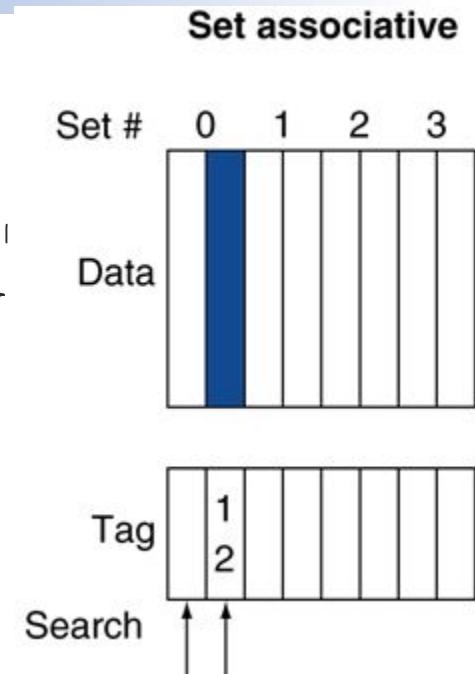
- A cache block can only go in one spot in the cache.
- It makes a cache block very easy to find
- but it's not very flexible about where to put the blocks.



2-Way Set Associative

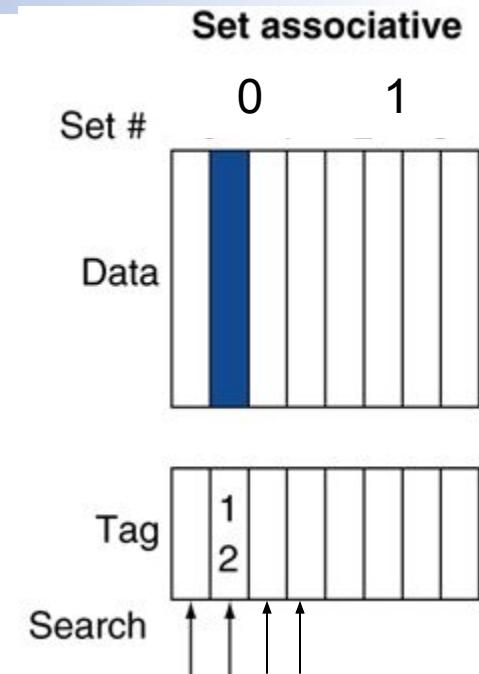


৫ blocks বাই set
অতি set টি ২টি
বরে entry রাখা যাব।
কেবল রাখে এটা
random.



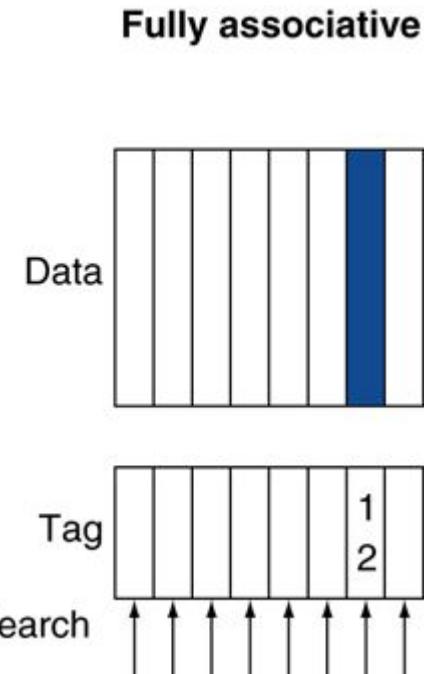
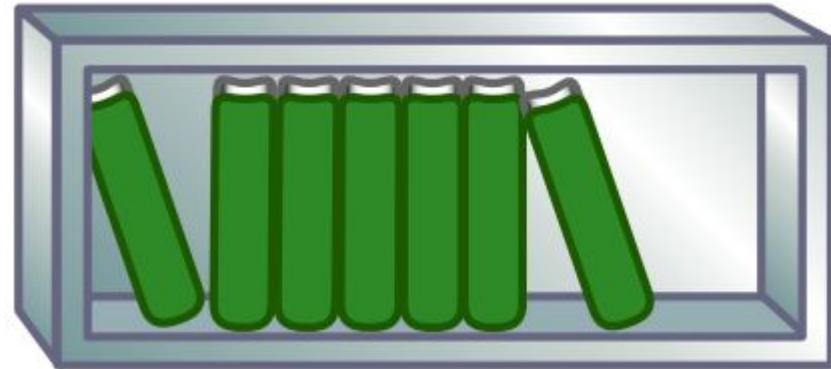
- This cache is made up of sets that can fit two blocks each.
- The index is now used to find the set
- The tag helps find the block within the set.

4-Way Set Associative



- Each set here fits four blocks,
- So there are fewer sets.
- As such, fewer index bits are needed.

Fully Associative



- No index is needed, since a cache block can go anywhere in the cache.
- Every tag must be compared when finding a block in the cache
- but block placement is very flexible!

Spectrum of Associativity

- For a cache with 8 entries

**One-way set associative
(direct mapped)**

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

for set 2 blocks

Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

· 4 entries
per set 1

Eight-way set associative (fully associative)

Tag	Data												

8 entries
1 set 1/8



Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

as পুরুষ রেন্ডা মাই

0 modulo 4 = 0	8 modulo 4 = 0	0 modulo 4 = 0	6 modulo 4 = 2	8 modulo 4 = 0	Block address	Cache index	Hit/miss	Cache content after access			
					0	1	2	3			
					0	0	miss	Mem[0]			
					8	0	miss	Mem[8]			
					0	0	miss	Mem[0]			
					6	2	miss	Mem[0]			
					8	0	miss	Mem[8]		Mem[6]	Mem[6]

Associativity Example

Compare 4-block caches

- Direct mapped, 2-way set associative, fully associative
- Block access sequence: 0, 8, 0, 6, 8

replacement scheme :

least recently used to replace.
recently accessed to evict.

2-way set associative

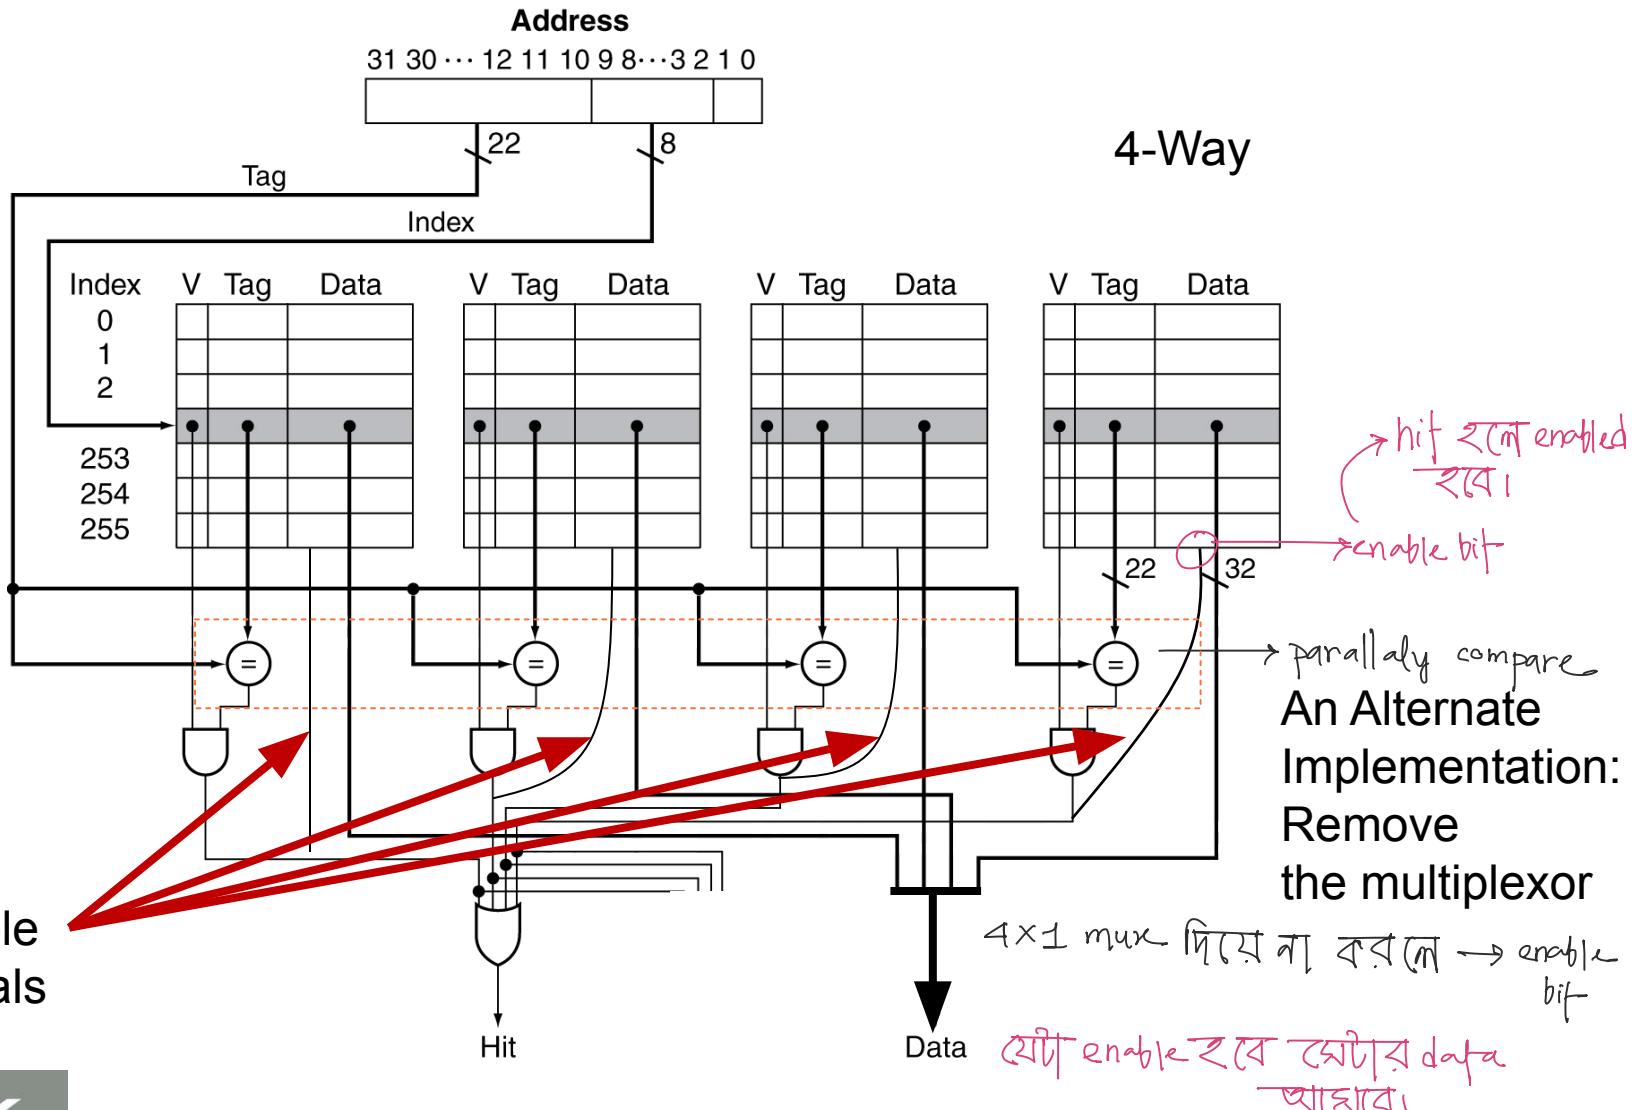
Block address	Cache index	Hit/miss	Cache content after access	
			- Set 0	Set 1
0 modulo 2 = 0	0	miss	Mem[0]	
8 modulo 2 = 0	8	miss	Mem[0]	Mem[8]
0 modulo 2 = 0	0	hit	Mem[0]	Mem[8]
6 modulo 2 = 0	0	miss	Mem[6]	Mem[6]
8 modulo 2 = 0	8	miss	Mem[8]	Mem[6]

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, **fully associative**
 - Block access sequence: 0, 8, 0, 6, 8
- Fully associative

Block address		Hit/miss	Cache content after access		
0		miss	Mem[0]		
8		miss	Mem[0]	Mem[8]	
0		hit	Mem[0]	Mem[8]	
6		miss	Mem[0]	Mem[8]	Mem[6]
8		hit	Mem[0]	Mem[8]	Mem[6]

Set Associative Cache Organization



Replacement Policy

- Direct mapped: no choice
- Set associative
 - Prefer non-valid entry, if there is one
 - Otherwise, choose among entries in the set
- Least-recently used (LRU)
 - Choose the one unused for the longest time
 - Simple for 2-way, manageable for 4-way, too hard beyond that
- Random
 - Gives approximately the same performance as LRU for high associativity

n -way set. n রেটখ্যাল চোস্চ

বোল্টি যদির আগে use হয় track করা লাগব।

2 way হলে \rightarrow 1টি extra bit রাখা লাগব।

4 way \rightarrow 2টি bit লাগব।

> 4 হলে \rightarrow problem.



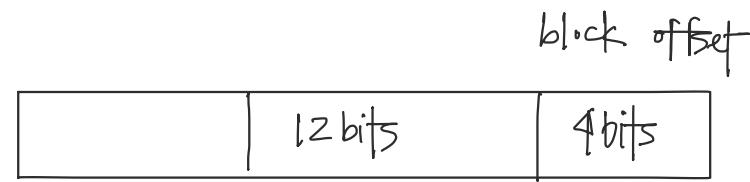
Tags versus Set Associativity

- Cache of 4096 blocks
- a 4-word block size
- 32-bit address,
- Find the total number of sets and the total number of tag bits for caches that are
 - direct mapped
 - two-way set associative
 - four-way set associative
 - fully associative.



Direct mapped:

$$2^{12} = 4096$$



$$\# \text{ sets} = 4096$$

$$32 - 12 - 2 - 2 = 16$$

↑
Tag bits
(per block)

$$\text{Total} = 16 \times 4096$$

2 way Set Associative:

$$\text{index} = 11$$

$$\# \text{ sets} = 2048$$

$$32 - 11 - 2 - 2 = 17$$

4 way Set Associative:

$$32 - 10 - 2 - 2 = 18$$

$$\# \text{ sets} = 1024$$

Fully Set Associative:

$$32 - 0 - 2 - 2 = 28$$

$$\# \text{ sets} = 1$$

Tags versus Set Associativity

- 4096 blocks; 32 bit address
- 4 words block = $16 (2^4)$ Bytes block
- So for tag + index $32 - 4 = 28$ bits



- Direct Mapped:
 - $4096 (2^{12})$ 1-way set => 12 bit index
 - 16 bit tag (28 - 12)
 - $4096 \text{ entries} \times 16 \text{ bit tag} = 64 \text{ K bits tag}$

Tags versus Set Associativity

- 4096 blocks; 32 bit address
- 4 words block = $16 (2^4)$ Bytes block
- So for tag + index $32 - 4 = 28$ bits



- 2-way set associative :
 - $4096/2 = 2048 (2^{11})$ sets => 11 bit index
 - 17 bit tag ($28 - 11$)
 - $4096 \text{ entries} \times 17 \text{ bit tag} = 68 \text{ K bits tag}$

Associativity $\uparrow \rightarrow$ Tag bit $\uparrow \rightarrow$ overhead $\uparrow \rightarrow$ costing \uparrow

Tags versus Set Associativity

- 4096 blocks; 32 bit address
- 4 words block = $16 (2^4)$ Bytes block
- So for tag + index $32 - 4 = 28$ bits



- 4-way set associative :
 - $4096/4 = 1024 (2^{10})$ sets => 10 bit index
 - 18 bit tag (28 -10)
 - $4096 \text{ entries} \times 18 \text{ bit tag} = 72 \text{ K bits tag}$

Tags versus Set Associativity

- 4096 blocks; 32 bit address
- 4 words block = $16 (2^4)$ Bytes block
- So for tag + index $32 - 4 = 28$ bits



- Fully set associative :
 - No index
 - 28 bits tag (28 - 0)
 - $4096 \text{ entries} \times 28 \text{ bit tag} = 112 \text{ K bits tag}$

Tags versus Set Associativity

- 4096 blocks; 4-words block; 32-bit address
- 4-words block = $16 (2^4)$ Bytes block
- So, for tag + index has $32 - 4 = 28$ bits.

28 bits (tag + Index)		4 bits
Associativity	Tag bits (K)	
Direct mapped (1 Way)	64	
2 Way Set Associative	68	
4 Way Set Associative	72	
Fully Associative	112	

costing almost
2x 2^{12}
(bit বাড়ালে)

Multilevel Caches

- Primary cache attached to CPU
 - Small, but fast
- Level-2 cache services misses from primary cache
 - Larger, slower, but still faster than main memory
- Main memory services L-2 cache misses
- Some high-end systems include L-3 cache



Multilevel Cache Example

Given

- CPU base CPI = 1, clock rate = 4GHz
- Miss rate/instruction @ primary cache = 2%
- Main memory access time = 100ns
- 4 GHz => 0.25ns cycle length

টাই 2% প্রিমিয়া
memory stall হবে।

ns to clock cycle
 $100/0.25$

With just primary cache

- Miss penalty = $100\text{ns}/0.25\text{ns} = 400$ cycles
- Effective CPI = $1 + 0.02 \times 400 = 9$

per instruction
($\frac{\square}{I}$)

$$1 + 0.02 \times 400 \times 1$$

total instruction 1 রেট



Example (cont.)

- Now add L-2 cache

- Access time = 5ns

primary miss
level 2 hit

CPU base CPI = 1,
clock rate = 4GHz

Global miss rate to main memory = 0.5%

- Primary miss with L-2 hit

- Penalty = $5\text{ns}/0.25\text{ns} = 20 \text{ cycles}$

- primary Miss
rate = 2%
- Primary miss with L-2 miss

$$\text{CPI} = 1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$$

- Performance ratio = $9/3.4 = \underline{\underline{2.6}}$

miss rate → করছে
performance → ডাকা হবে

→ যদি global
context এ তাহু-
add করতে পারিব।
→ Local হলে মিথ্যে
দেখা লাগবে

Multilevel Cache Considerations

■ Primary cache

- Focus on minimal hit time

primary
cache → minimal hit time (target)

■ L-2 cache

- Focus on low miss rate to avoid main memory access
- Hit time has less overall impact

L2 → miss rate কমানো
target করা নি miss হলে
mem G ঘটে হবে।

■ Results

- L-1 cache usually smaller than a single cache
- L-1 block size smaller than L-2 block size



Interactions with Advanced CPUs

- Out-of-order CPUs can execute instructions during cache miss
 - Pending store stays in load/store unit
 - Dependent instructions wait in reservation stations
 - Independent instructions continue
- Effect of miss depends on program data flow
 - Much harder to analyse
 - Use system simulation

→ in order execution

ରହନା।

— cache miss ହୁଏ

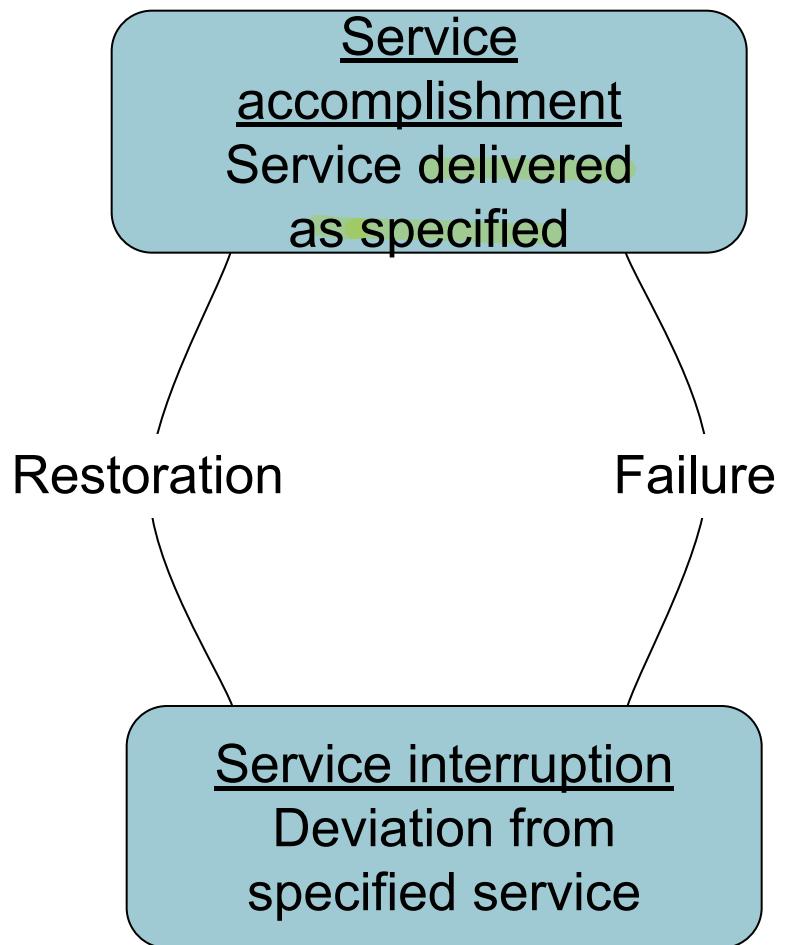
— ଅନ୍ୟ ଯେଉଁ ins.

ପର୍ଯ୍ୟନ୍ତ ଉପର

dependent ରୀତିରେ

exec.

Dependability



- Fault: failure of a component
 - May or may not lead to system failure

Dependability Measures

- Reliability: mean time to failure (MTTF) → বাত যন্ত্র মুশি
র ব্যবহাৰ
- Service interruption: mean time to repair (MTTR)
- Mean time between failures
 - $MTBF = MTTF + MTTR$ → ধৰণীক সেইন্টুৰুৰ কৰণৰ সময়
- Availability = $MTTF / (MTTF + MTTR)$
- Improving Availability
 - Increase MTTF: fault avoidance, fault tolerance, fault forecasting
→ ২টা পেৰু দয়া: একটা পিঁ
কৰলো অন্যটা service কৰিছো
 - Reduce MTTR: improved tools and processes for diagnosis and repair

Nines of Availability

- We want availability to be very high.
- One shorthand is to quote the number of “nines of availability” per year.
 -

# of Nines	% of uptime
One	90
Two	99
Three	99.9
Four	99.99
Five	99.999

Nines of Availability

- Given 365 days per year, which is $365 * 24 * 60 = 525,600$ minutes.
- Then the shorthand is decoded as follows:
- $90\% \Rightarrow 525,600 * 0.1$ downtime = 52560 minutes = $52560/(60*24) = 36.5$ days.

# of Nines	% of uptime	Downtime/Year
One	90	36.5 days
Two	99	3.65 days
Three	99.9	526 min
Four	99.99	52.6 min
Five	99.999	5.26 min

MTTF and AFR

- MTTF is a reliability measure.
- A related term is *annual failure rate* (AFR)
 - The percentage of devices that would be expected to fail in a year for a given MTTF.
 - Hours in a year / MTTF
 - When MTTF gets large it can be misleading
 - while AFR leads to better intuition

High MTTF and AFR

- Some disks today are quoted to have a 1,000,000-hour MTTF.
 - $1,000,000 / (365 * 24) = 114$ years
 - they practically never fail???
- Warehouse scale computers that run Internet services such as Search might have 50,000 servers.
- Assume each server has 2 disks.
- How many disks we would expect to fail per year.



High MTTF and AFR

- One year => $365 * 24 = 8760$ hours.
- A 1,000,000-hour MTTF means an AFR of $8760/1,000,000 = 0.876\%$.
- We have $50000 * 2 = 100,000$ disks
- we would expect $0.00876 * 100,000 = 876$ disks to fail per year
- On average more than $(876/365)$ 2 disk
failures per day.
tolerable

Question:

- Cache design
- math

✓
CT-SYII

Virtual Memory

প্রতিটি program এর নিজস্ব address space
যাববে, অটোর যেখনে ল্যাণ্ড মে write
করতে পারবে।

- Use main memory as a “cache” for secondary (disk) storage
 - Managed **jointly** by CPU hardware and the operating system (OS)
- Programs share main memory
 - Each gets a private virtual address space holding its frequently used code and data
 - Protected from other programs
- CPU and OS translate virtual addresses to physical addresses
 - VM “block” is called a page
 - unit of data hard disk থেকে
পচে RAM ৰ বাখ্যে
 - VM translation “miss” is called a page fault

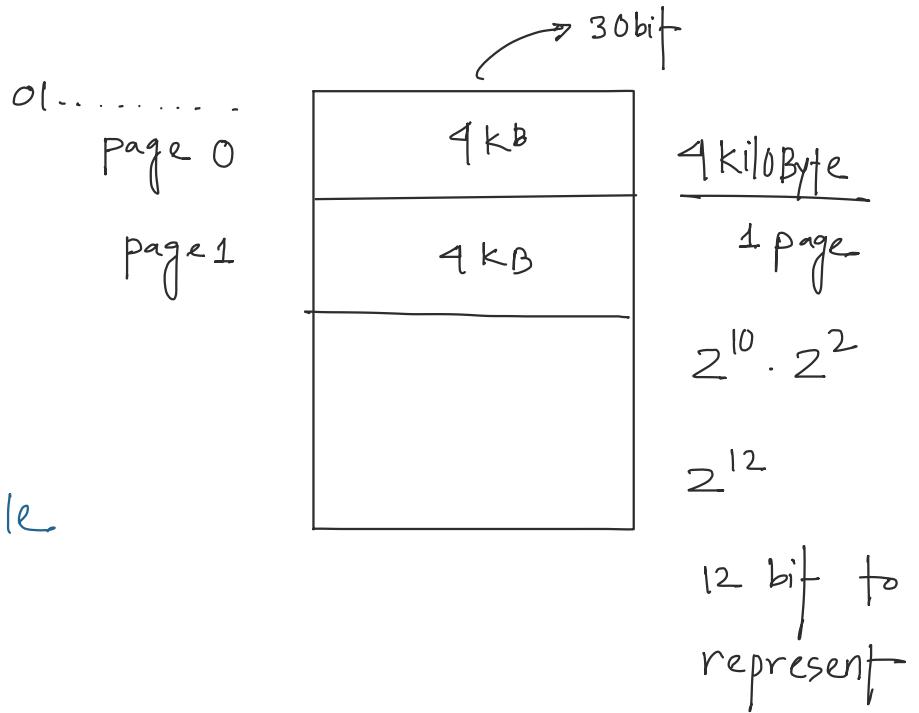
virtual address থেকে

physical address এ map করার

জন্য page table লাগে।

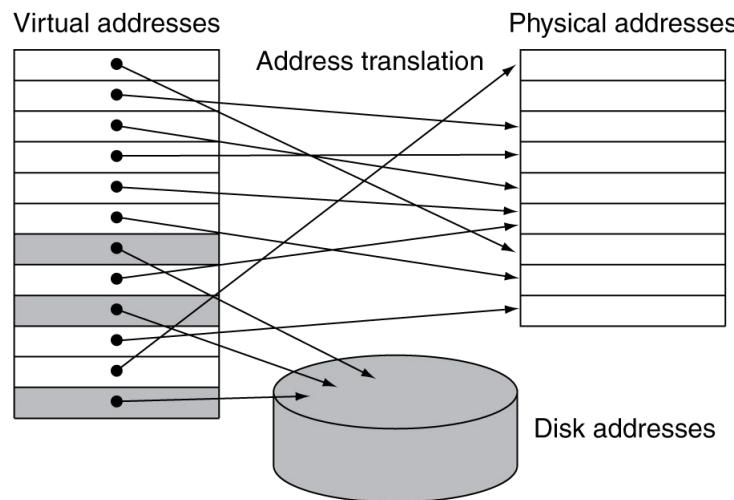
এতো process এর জন্য আলাদা page table

নিলে \rightarrow no. of entries ঘনের (বড়ে) যায়।

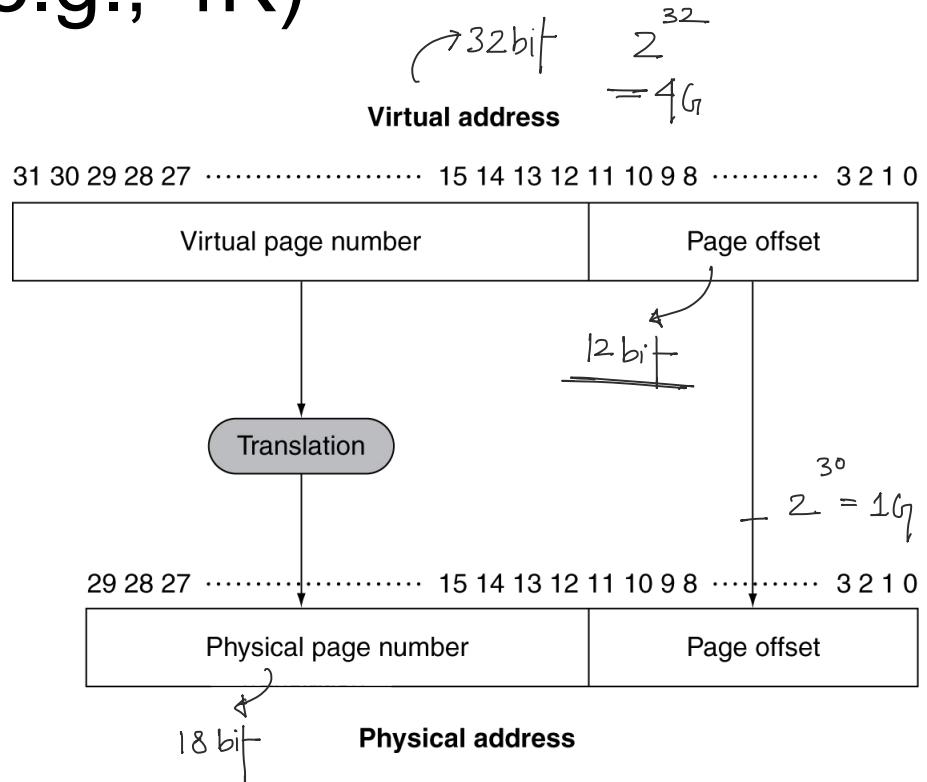


Address Translation

■ Fixed-size pages (e.g., 4K)



design focus → minimal page fault



OS protection দ্বারা যেন একটি program এর memory পার্শ থালে program access না হোলো।

Page Fault Penalty

- On page fault, the page must be fetched from disk
 - Takes millions of clock cycles
 - Handled by OS code
- Try to minimize page fault rate
 - Fully associative placement
 - But this needs costly search!!!
 - Smart replacement algorithms

Page Tables

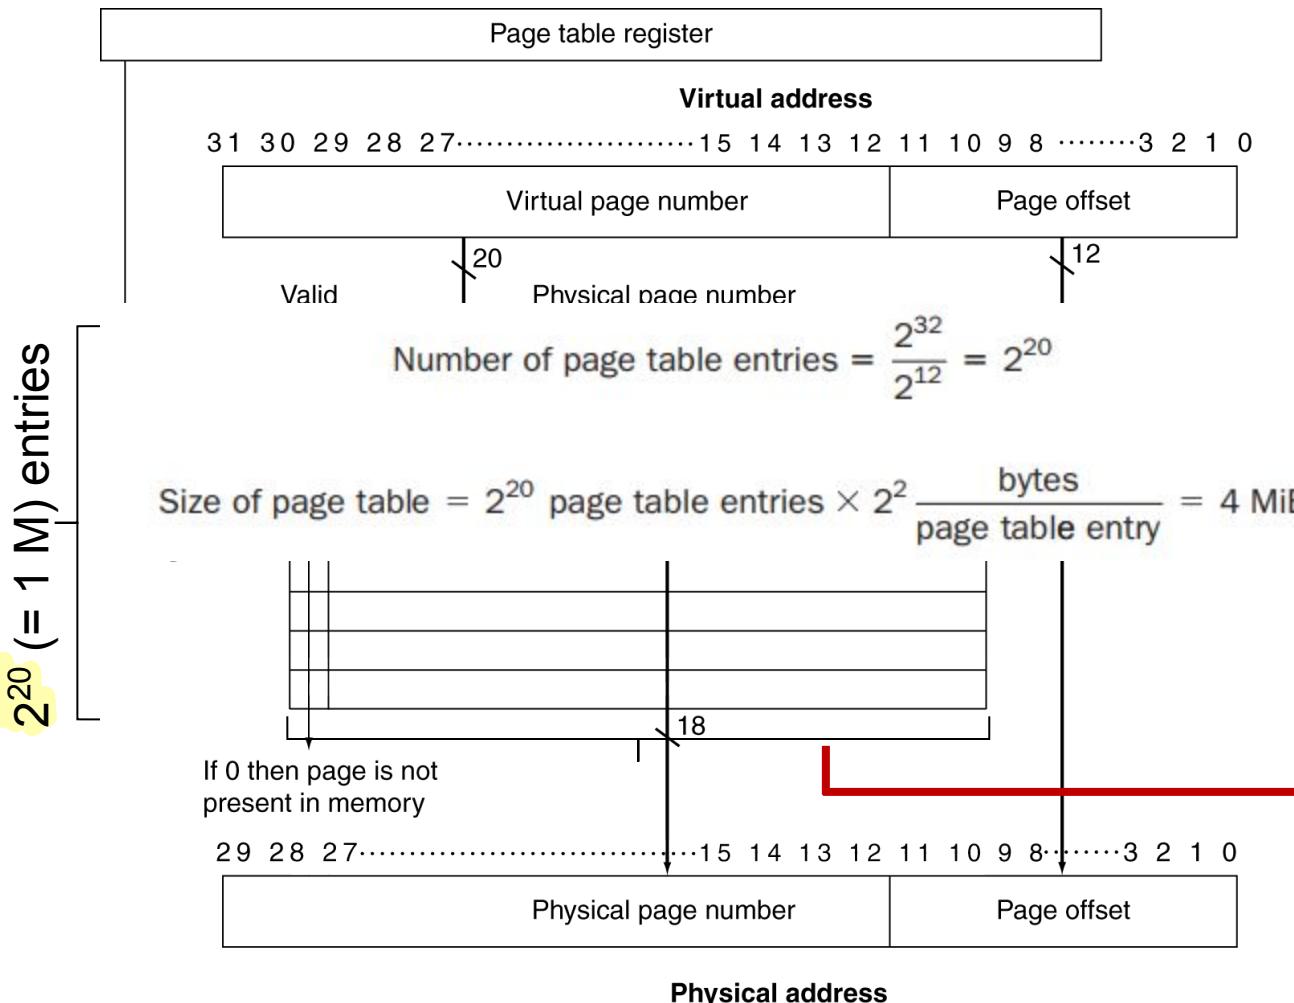
এতিঁ process এর জন্য একটি করে page table
যাবে। \hookrightarrow নিচের virtual space থেকে

- Stores placement information
 - Array of page table entries, indexed by virtual page number
 - Page table register in CPU points to page table in physical memory
- If page is present in memory
 - PTE stores the physical page number
 - Plus other status bits (referenced, dirty, ...)
- If page is not present
 - PTE can refer to location in swap space on disk



valid bit → physically map করা আছে বিনা।

Translation Using a Page Table



What is
the size
of the
page
table?

19 bits;
but 32 bits
wide usually

4 bytes $\times 2^{20}$ entries

4 MB \rightarrow size of
a page table

Page Table Size Issues

- Size of page table => 4 MB (per process)
- What about 100 processes
 - Each with its own page table?
- What will happen if we have 64-bit addresses (according to prev. calc.)?
 - $2^{(64 - 12)} = 2^{52}$ entries!!!



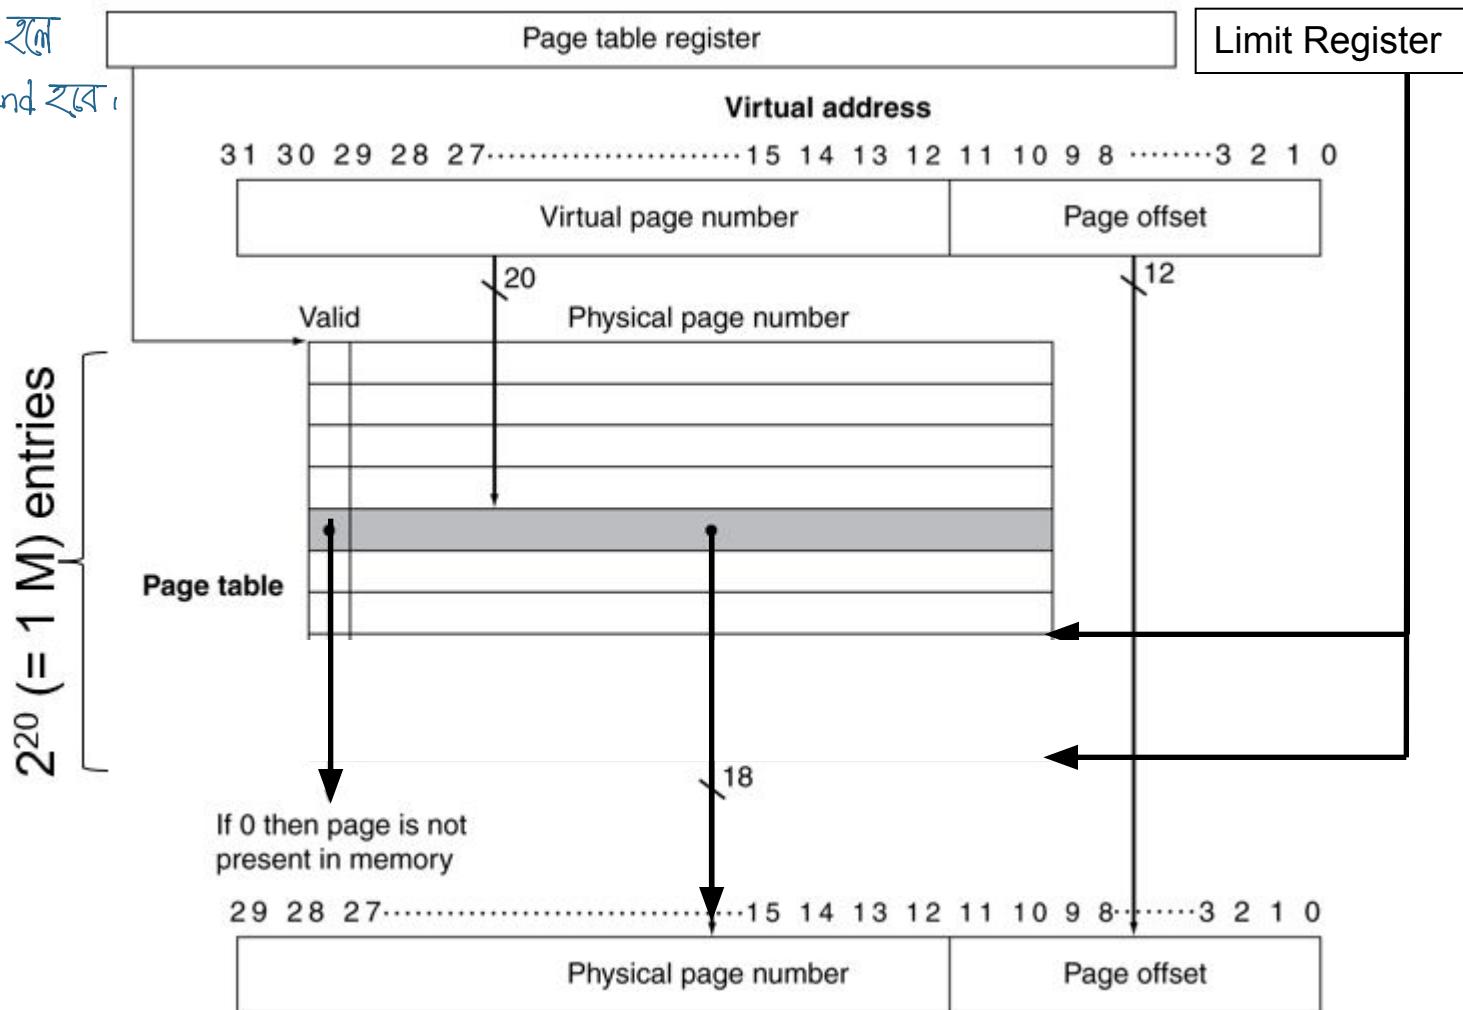
- There are techniques to reduce the amount of storage required for the page table.

Techniques: Limit Registers

- Keep a limit register that restricts the size of the page table for a given process.
 - If the virtual page number becomes larger than the contents of the limit register, entries must be added to the page table.
 - This technique allows the page table to grow as a process consumes more space.
 - Thus, the page table will only be large if the process is using many pages of virtual address space.
 - This technique requires that the address space expand in only one direction.

Techniques: Limit Registers

higher order এর
পুরোটা extend হবে।

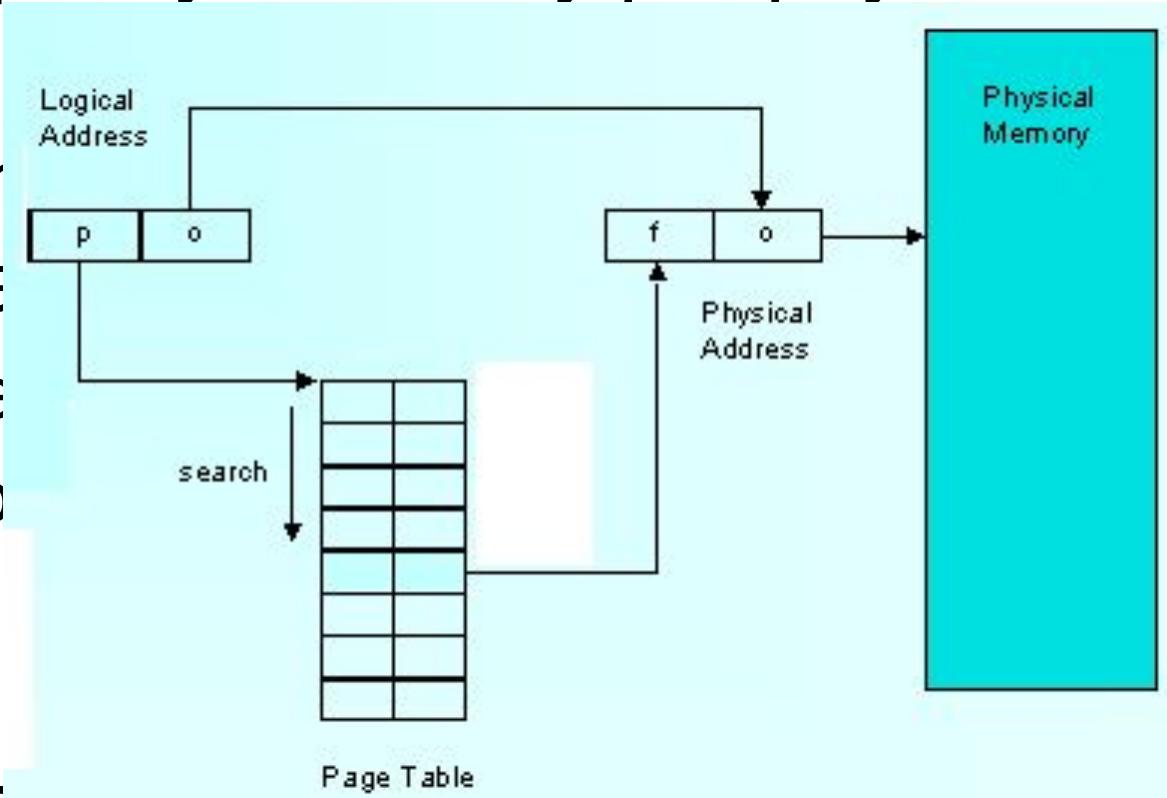


Techniques: Two Limits

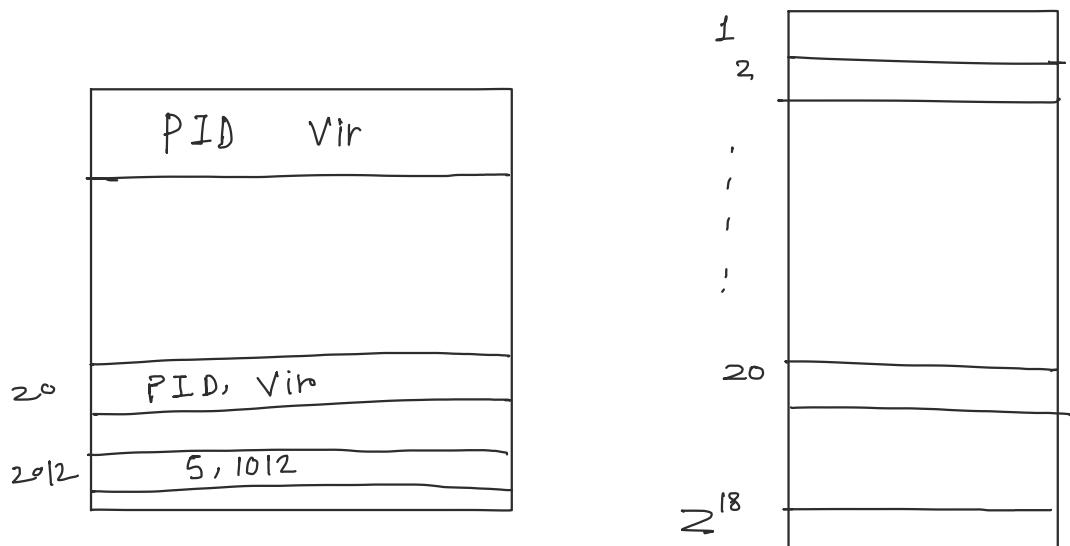
- Two separate page tables and two separate limits.
- | | |
|-------|--|
| Stack | <ul style="list-style-type: none">■ One grows from highest address down |
| Heap | <ul style="list-style-type: none">■ One grows from lowest address up■ So, address space is divided into 2 segments.■ High-order bit of an address usually determines which segment<ul style="list-style-type: none">■ i.e., which page table to use for that address.■ So, each segment can be as large as one-half of the address space.■ A limit register for each segment specifies the current size of the segment, which grows in units of pages. |

Techniques: Inverted Page Table

- Keep only one entry per physical block (i.e., such that each page table entry corresponds to one physical page)
- Such that each page table entry corresponds to one physical page.
- we can do this by using an inverted page table.
- So, we can do this by using an inverted page table.
- May be faster because it reduces the number of lookups required to find a page.
- To make the lookups faster



all process এর জন্য একটি page-table maintain — inverted page table



$$2^{18} \times 4 = 2^{22}$$

প্রতিটি entry 4 KB

PID \rightarrow process identifier

Vir \rightarrow virtual address

2012 page 6 5 no process, 1012 virtual address

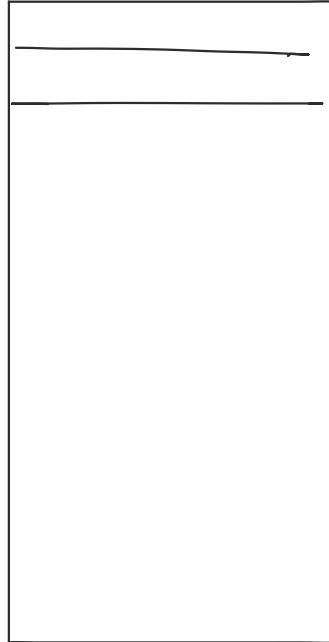
5 থেকে থেকে then check করাবে vir 1012 বিলক্ষণ index থেকে page no.

(physical)

problem: মুয়োটি search করা লাগে।

1MB দিয়েই যথগুলো process এর জন্যই page-table বানাতে পারিব।

V.A



2^o

111 ...

0...10
---|1

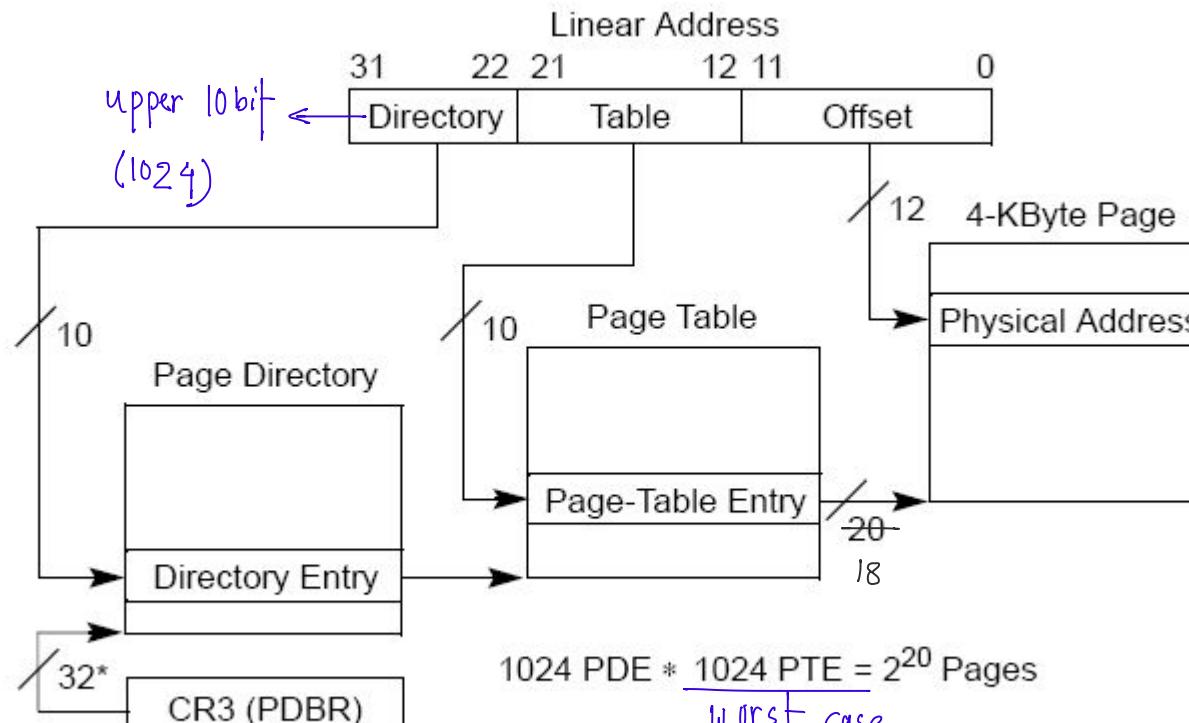
অতি V.A ব্যবহার করুন।

Techniques: Multiple Levels

- The first level maps large fixed-size blocks of virtual address space **Sometimes called segment**
- Each entry in the segment table:
 - indicates whether any pages in that segment are allocated
 - if so, points to a page table for that segment.
- Address translation happens:
 - by first looking in the segment table, using the highest-order bits of the address.
 - If the segment address is valid, the next set of high-order bits is used to index the page table



Techniques: Multiple Levels

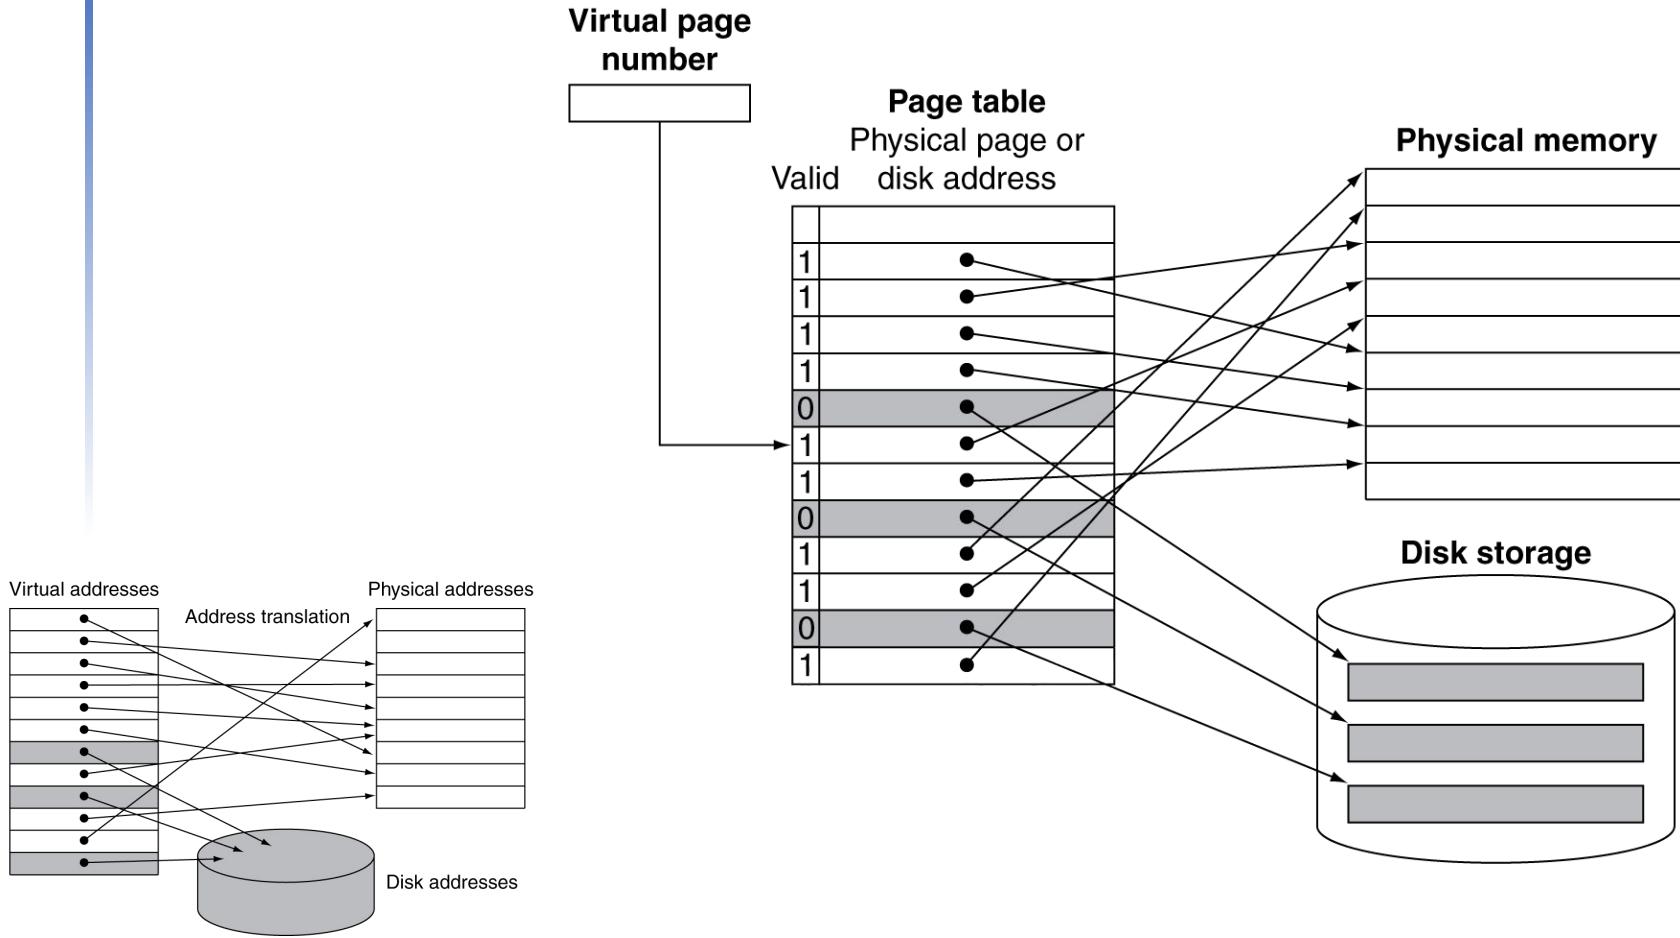


cache-memory থেকে \rightarrow block

memory - hard disk থেকে \rightarrow page
গড়ে আনে

average case গু প্রয়োবেন,

Mapping Pages to Storage



Page Fault

କୋଣେ ଏହାଟି page ଯାହିଁ memory ରୁ ନଥାଏ — page fault
valid bit ୦ ଥାଏ ।
— OS ରେ control ଦିଇଯାଇବା

- If the valid bit for a virtual page is off, a page fault occurs.
- The operating system must be given control.
 - This transfer is done with the exception mechanism processor LR
 - OS must find the page in the next level of the hierarchy harddisk
 - and decide where to place the requested page in main memory.

Swap Space

- Virtual addr. alone does not immediately tell us where the page is on disk.
 - OS usually creates the space on flash memory/disk for all the pages of a process
 - OS creates a data structure to record where each virtual page is stored on disk.
 - may be part of the page table
 - or an auxiliary data structure indexed in the same way as the page table
- OS also creates a data structure that tracks
- which processes and which virtual addresses use each physical page
- virtual address
এর mapping
থাকে harddisk
and physical
address এর
যাতে*

Replacement and Writes

- To reduce page fault rate, prefer least-recently used (LRU) replacement
 - Reference bit (aka use bit) in PTE set to 1 on access to page *use bit = 1 (recently accessed)*
 - Periodically cleared to 0 by OS
 - A page with reference bit = 0 has not been used recently
- Disk writes take millions of cycles
 - Block at once, not individual locations
 - Write through is impractical — ram and hard disk এ প্রক্ষার্তে update (use কৰা হয়না)
 - Use write-back — dirty bit লাগব।
 - Dirty bit in PTE set when page is written — আগের page

hard disk এ রেখে বস্তন page দিয়ে আবাত হবে।

যদি আগের dirty bit 1
খাতে

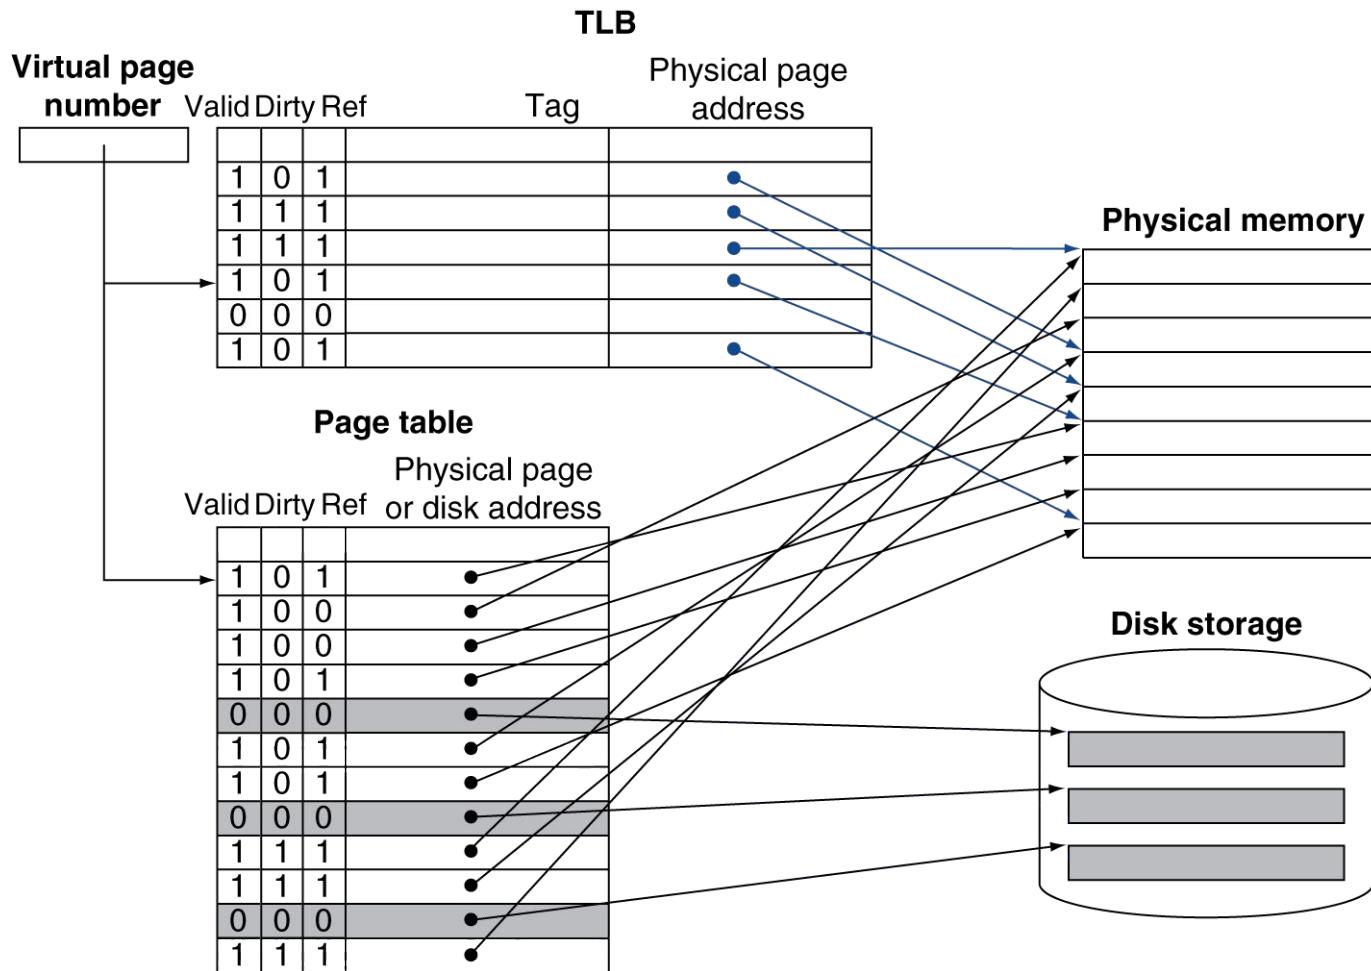
Fast Translation Using a TLB

- Address translation would appear to require extra memory references
 - One to access the PTE
 - Then the actual memory access
- But access to page tables has good locality
 - So use a fast cache of PTEs within the CPU
 - Called a Translation Look-aside Buffer (TLB)
 - Typical: 16–512 PTEs
 - Misses could be handled by hardware or software

TLB miss → corresponding page ram/hard disk
— থাকতে পারে।



Fast Translation Using a TLB



TLB Misses

- If page is in memory
 - Load the PTE from memory and retry
 - Could be handled in hardware
 - Can get complex for more complicated page table structures
 - Or in software
 - Raise a special exception, with optimized handler
- If page is not in memory (page fault)
 - OS handles fetching the page and updating the page table
 - Then restart the faulting instruction



TLB Miss Handler

- TLB miss indicates
 - Page present, but PTE not in TLB
 - Page not present **True Page Fault**
 - Handler copies PTE from memory to TLB
 - Then restarts instruction
 - If page not present, page fault will occur
 - The reference and dirty bits may change in TLB
 - So they must also be copied back to PTE for the TLB entry that is replaced **Write-Back Scheme**
- ram ৰ পৰি address টোয়ার কোর্সিং
page আছে
- write কৰাব
cache hierarchy → cache টোয়ার PT
PT টোয়ার harddisk



TLB: Associativity

Lower Miss Rate

- Some systems use small, fully associative TLBs
 - replacement choice becomes tricky
 - hardware LRU scheme is too expensive.
 - expensive software algorithm is also not feasible (unlike page faults)
 - Many systems provide some support for randomly choosing an entry to replace.
- Other systems use large TLBs, often with small associativity.

TLB misses
are much
more
frequent



Page Fault Handler

- Use faulting virtual address to find PTE
- Locate page on disk → OS এর বাছে mapping থাকে
- Choose page to replace
 - If dirty, write to disk first
- Read page into memory and update page table
- Make process runnable again
 - Restart from faulting instruction

Intrinsity FastMATH TLB

- 4 KiB (2^{12}) pages; 32-bit address space
 - So, virtual page number is $(32 - 12) = 20$ bits long
 - Physical address is same size as virtual address.
(32bit)
- TLB: 16 entries; fully associative
 - shared between the instruction and data
 - Each entry is 64 bits wide →
data cache ফিল্ড এবং
64 bit → 2 word
2 bit byte offset
1 bit word offset
... : 29 bit tag
 - a 20-bit tag (virtual page number for that TLB entry)
 - the corresponding physical page number (also 20 bits),
But
 - a valid bit, a dirty bit, and other bookkeeping bits.
 - Like most MIPS systems, it uses software to handle TLB misses.
*এখানে 12 bit page offset
20 bit (Indexing + tag)
TLB cache
ক্ষতি দা*
∴ fully associative → no index bit so 20 bit tag .

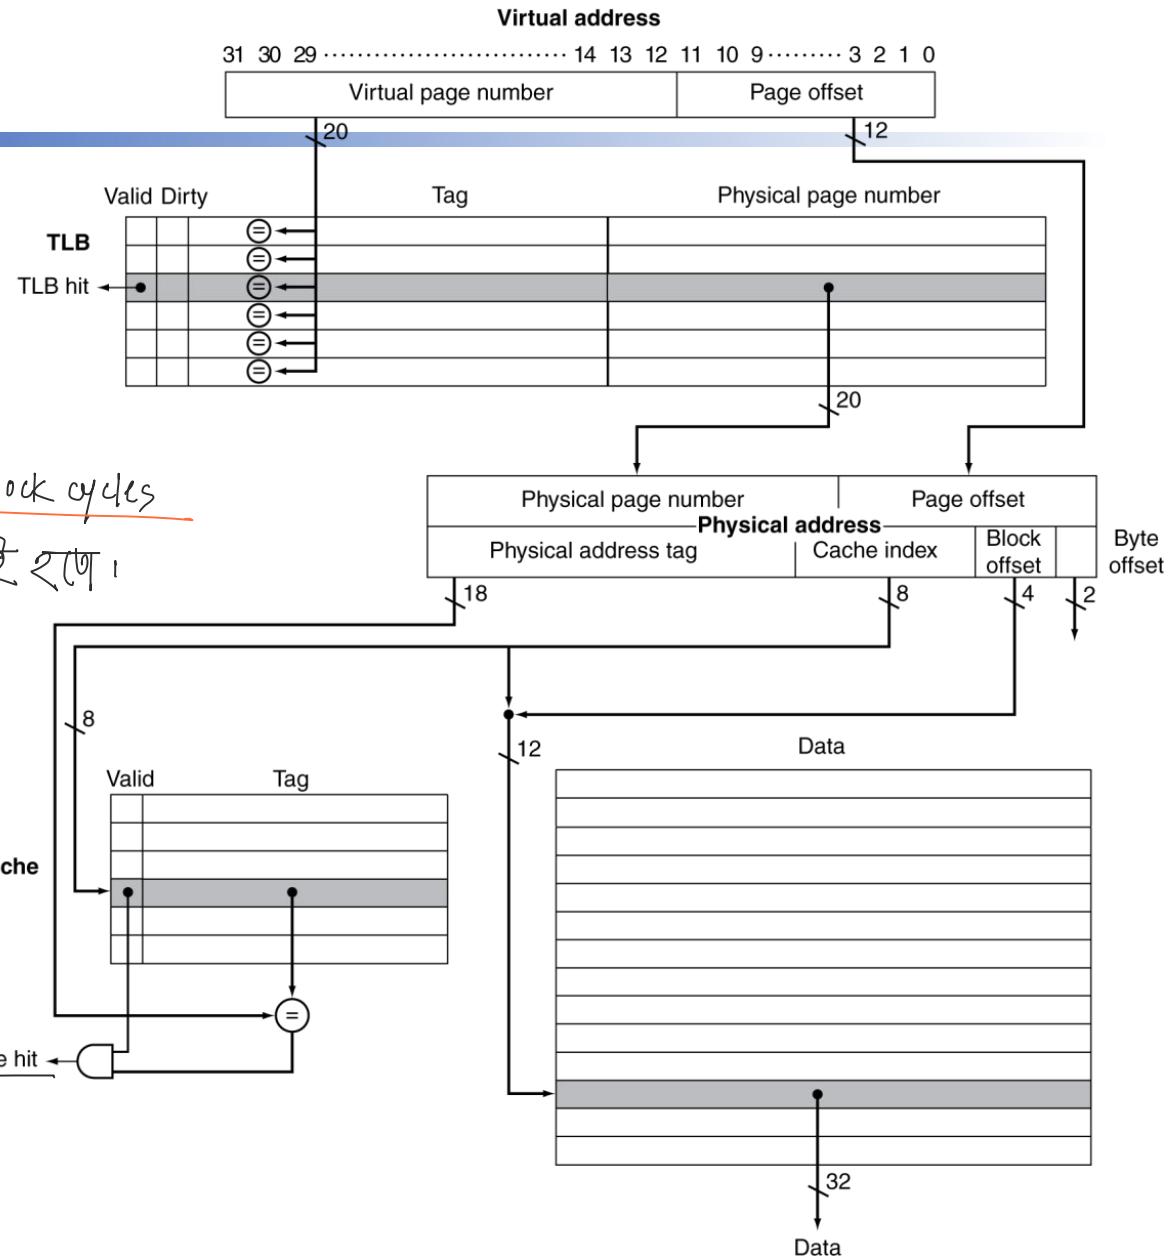
TLB and Cache Interaction

CAM Implementation

TLB Cache যদি না থাকলে \rightarrow
 PT দ্বাৰা প্ৰয়োজনীয়। Virtual
 address থেকে physical address
 পাওয়া হৈব (lw) + চাহত 100 clock cycles
 পাওয়াৰে, যদি 1/2 clock cycle হৈব হৈলে।
 \downarrow
 memory access কোমতো

2 bit byte offset
 4 bit word
 8 bit index
 ... tag 18 bit

Cache
 2^8 blocks \times 16 words/block



Content Addressable Memory (CAM)

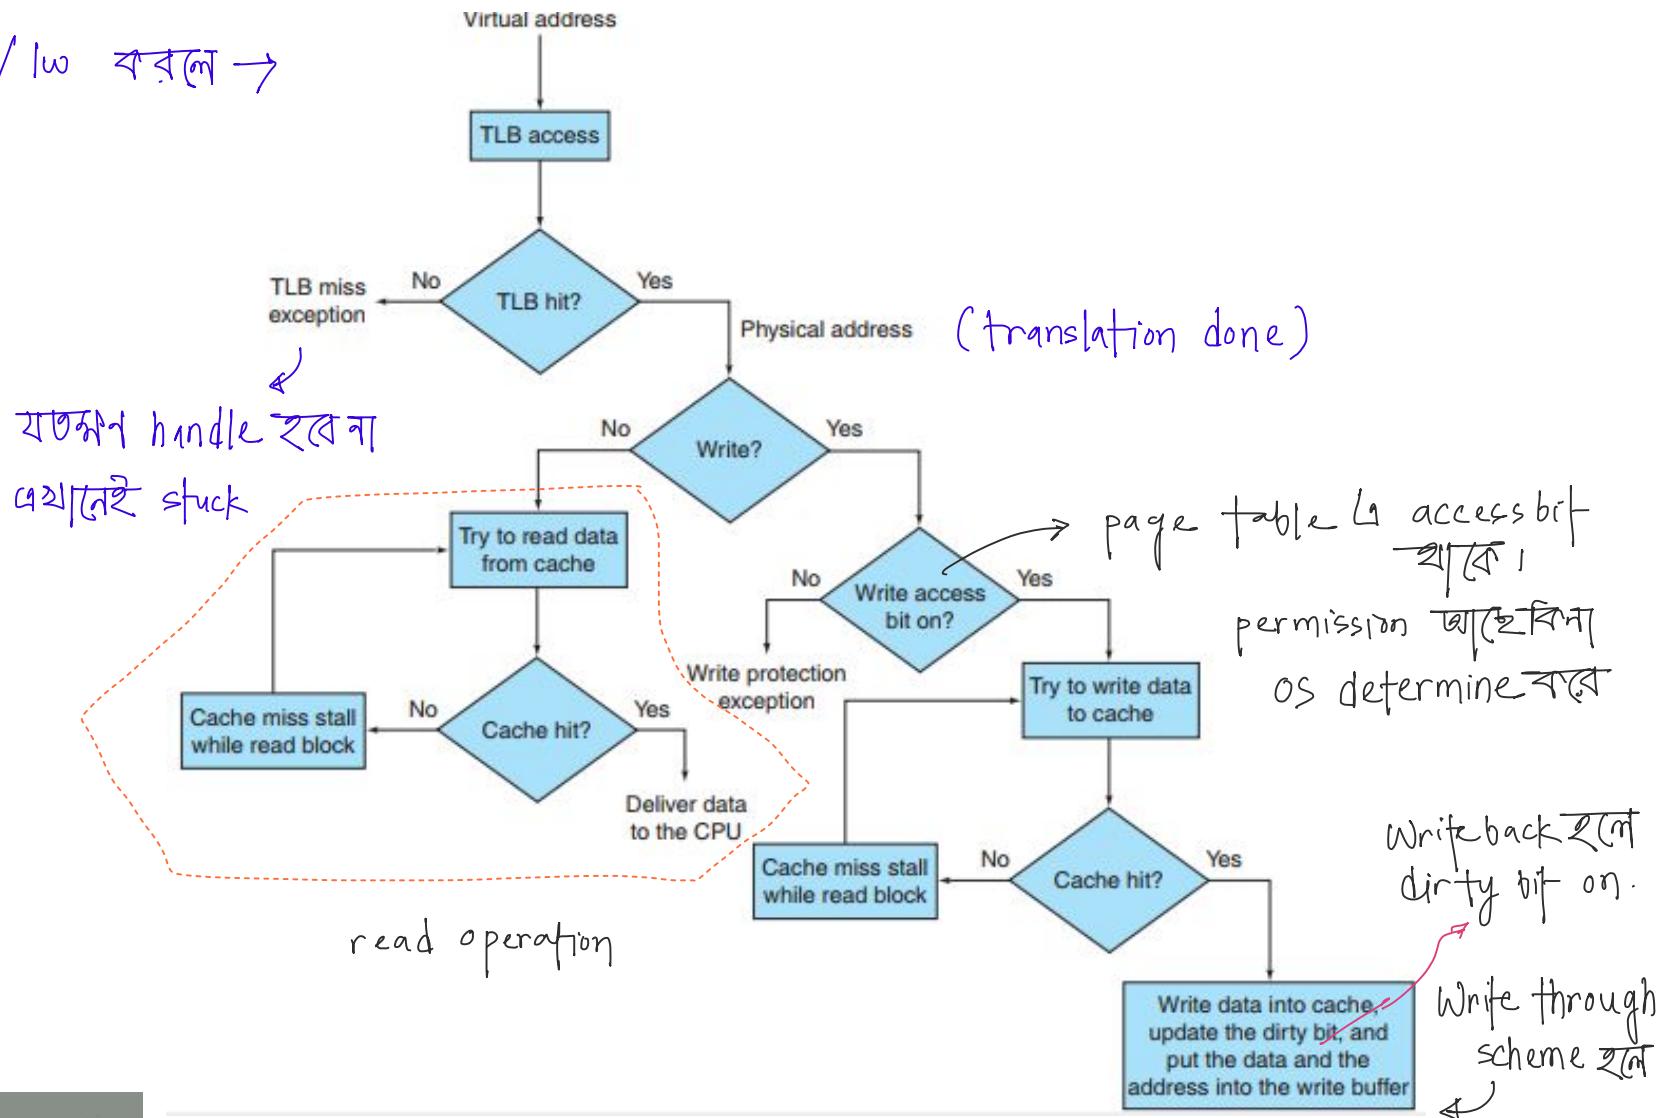
- CAM is a circuit that combines comparison and storage in a single device.
 - Does not supply an address and read a word like a RAM
 - you supply the data and the CAM looks to see if it has a copy and returns the index of the matching row.
 - With CAMs higher set associativity in cache can be implemented

RAM → address → value → return
CAM → value → " → address / " →
data → index



Processing a read or a write-through in the Intrinsity FastMATH TLB and cache

sw / lw করলে \rightarrow



TLB, Cache and VM Events Combined

Page ramରେ, mapping TLB କରିବାକୁ
but

ଏହା ପରେ data ଖୁଣ୍ଡଗେଣ୍ଟିଯିବା cache ଆନା ହେଲାକିମ୍ବା

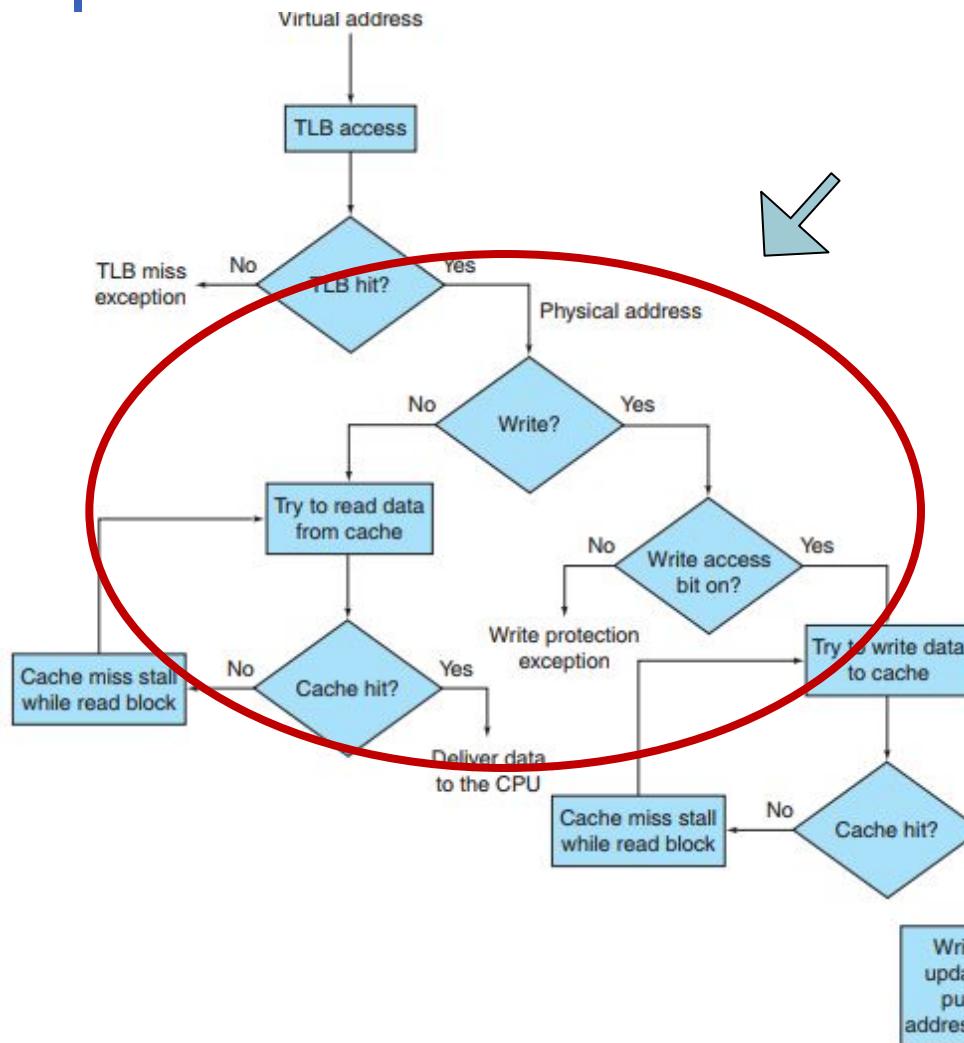
TLB	Page table	Cache	Possible? If so, under what circumstance?
Hit	Hit	Miss	Possible, although the page table is never really checked if TLB hits.
Miss	Hit	Hit	TLB misses, but entry found in page table; after retry, data is found in cache.
Miss	Hit	Miss	TLB misses, but entry found in page table; after retry, data misses in cache.
Miss	Miss	Miss	TLB misses and is followed by a page fault; after retry, data must miss in cache.
Hit	Miss	Miss	Impossible: cannot have a translation in TLB if page is not present in memory.
Hit	Miss	Hit	Impossible: cannot have a translation in TLB if page is not present in memory.
Miss	Miss	Hit	Impossible: data cannot be allowed in cache if the page is not in memory.

PT Hit → disk ଥିଲେ ram କିମ୍ବା ଅନା ହେଲେ ଯେ page ଖୁଣ୍ଡଗେଣ୍ଟି

TLB miss ହେଲେ → PT Hit check କରାଯାଇଲା

hit " → PT check କରାଯାଇଲା

Physically Addressed Cache



- Here, cache is “physically addressed” and “physically tagged”
- Time to access memory for a cache hit include:
 - TLB access time
 - Cache access time

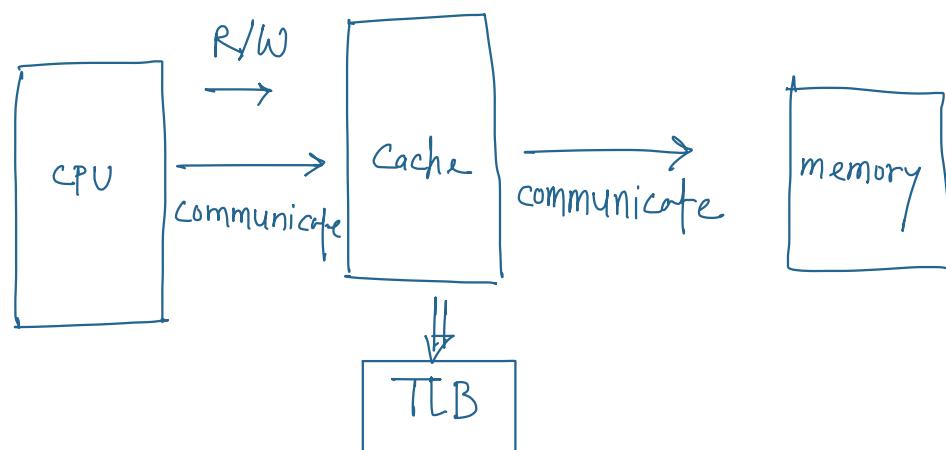
Latency বাটো
- Of course, these accesses can be pipelined.

Virtually Indexed Cache

cache index কোথায় থাকে virtual address ?
tag bit কোথায় থাকে virtual address ?

- Alternatively, the processor can index the cache with a virtual address VIVT
 - *Virtually indexed and Virtually tagged cache*
- Here, TLB is unused during the normal cache access
 - Reduce cache latency
- Cache miss=> the processor needs to translate the address to a physical address
 - so that it can fetch the cache block from main memory.

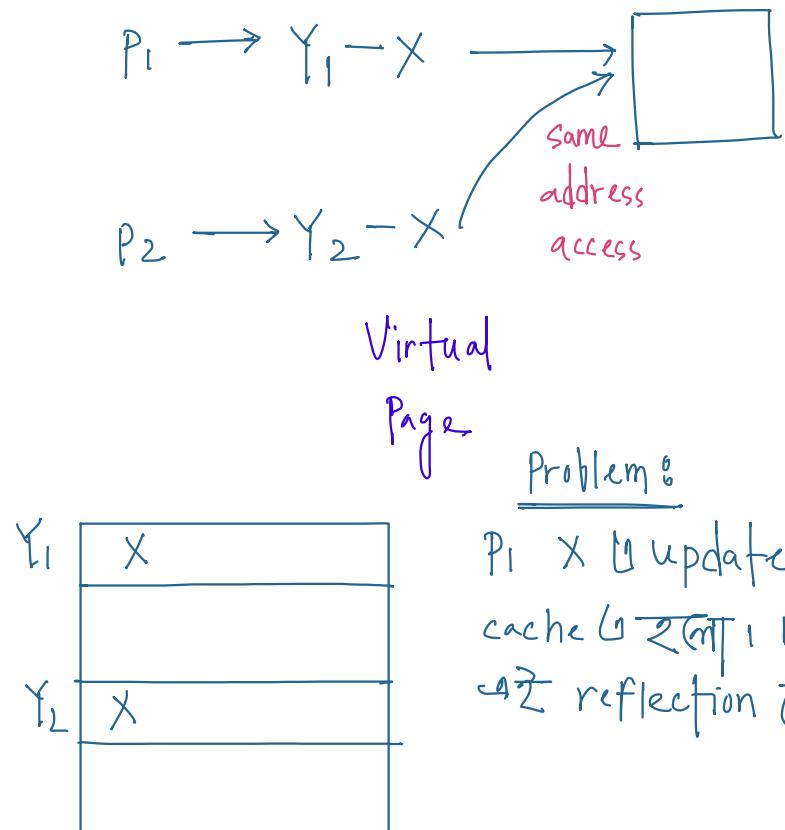




Cache ରିଟେଣ୍ଡିନ୍ଡିକେସନ୍

ନାଥାବଳେ mem

ଯେତେ ଏହି ପଢ଼େ ଯାଏନାମାରୁବୋ Physical address ଲାଗିବେ, ଆମାଦେଇ ବାହେ
virtual address ଯାଏନ୍ତି TLB ମୋଟି physical address ଏମଣ୍ଡ ଏବଂ



virtually indexed
 advantage — TLB লাগছে না
 physical address মিশে নাই
 বিকৃত cache miss রয়ে TLB access
 করা লাগবে and memory থেকে data
 আনা লাগবে।

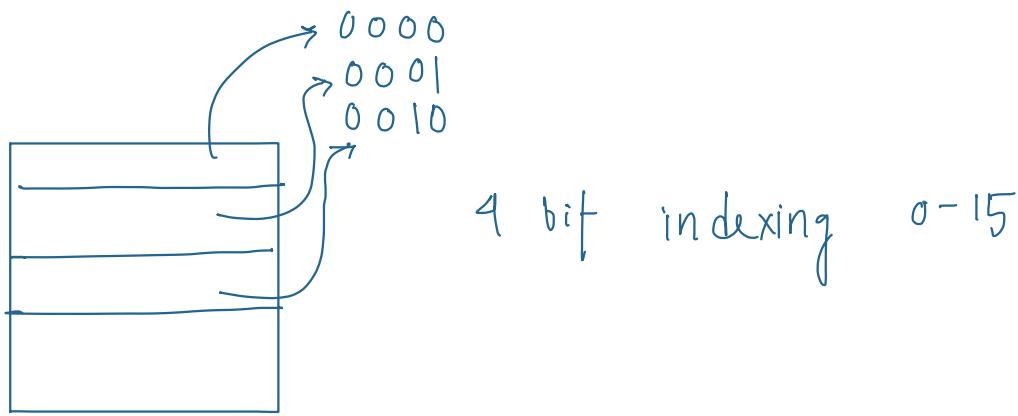
Soln: interprocess communication
 করতে দিয়োনা।

Aliasing in VIVT Cache

- Aliasing occurs when the same object has two names
 - Two virtual addresses for the same page.
 - May happen for shared pages between processes
 - This ambiguity creates a problem:
 - A word on such a page may be cached in two different locations, each corresponding to different virtual addresses.
 - One program may write the data without the other program being aware that the data had changed.
- Solution to Aliasing Issues
 - either introduce design limitations on the cache and TLB to reduce aliases
 - or require the operating system, and possibly the user, to take steps to ensure that aliases do not occur

Virtually Indexed Physically Tagged

- Physical Tag using just the page-offset portion of the address, which is really a physical address since it is not translated
- These designs, which are *virtually indexed but physically tagged*, attempt to achieve the performance advantages of virtually indexed caches with the architecturally simpler advantages of a physically addressed cache.
 - There is no alias problem in this case.

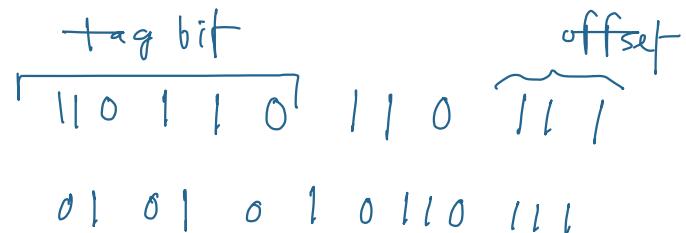


8 bit 0-255 পর্যন্ত

256

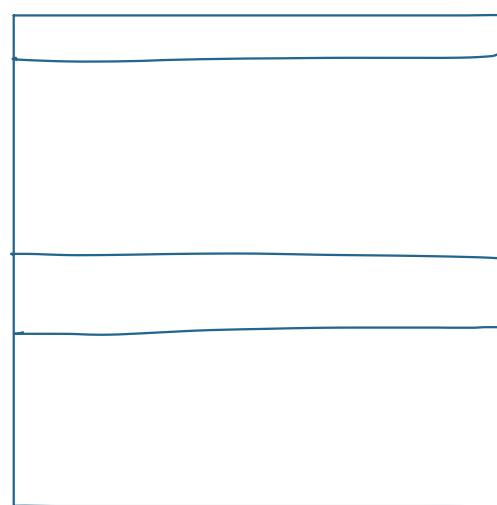
TLB access data cache

access — in parallel হয়



?

cache এর tag bit physical
address এর tag bit match করে
বিনা check.



index virtual address

থেকে tag physical " থেকে ,

তাই indexing পর্যন্ত TLB লাগে না

virtual address থেকে indexing

আমা আলাদা বোলো meaning দেই

indexing এর virtual address থেকে 8 bit দ্বারা physical

থেকে নেওয়া same! ফোর্মাট store ক্ষমতা ,

indexing করে যাওয়ার
পরে entry টি tag এর
সাথে compare করালাগে।

tag bit →

virtual address থেকে যে

physical address মাঝে

দেখির upper bits

TLB access করাবে

index virtual address

থেকে tag physical " থেকে ,

তাই indexing পর্যন্ত TLB লাগে না

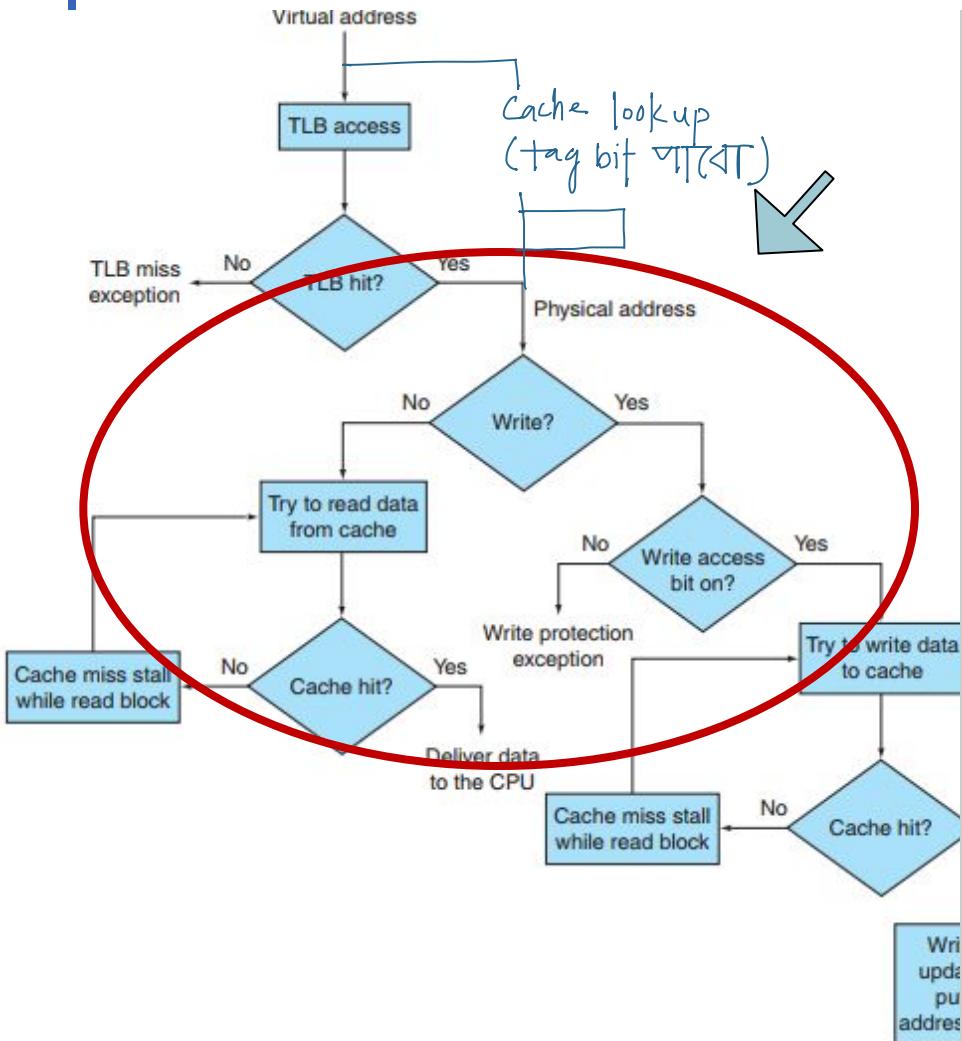
virtual address থেকে indexing

আমা আলাদা বোলো meaning দেই

indexing এর virtual address থেকে 8 bit দ্বারা physical

থেকে নেওয়া same! ফোর্মাট store ক্ষমতা ,

Physically Addressed Cache



- Here, cache is “physically addressed” and “physically tagged”
 - Time to access memory for a cache hit include:
 - TLB access time
 - Cache access time
 - Of course, these accesses can be pipelined.
- Latency
ব্যাপ্তি*

Memory Protection

- Different tasks can share parts of their virtual address spaces
 - But need to protect against errant access
 - Requires OS assistance
- Hardware support for OS protection
 - Privileged supervisor mode (aka kernel mode)
OS এই switching handle করে।
 - Privileged instructions
 - Page tables and other state information only accessible in supervisor mode
 - System call exception (e.g., syscall in MIPS)

e.g. page fault

operating system এর বাছে যেটা হলো যাবে

Memory Protection: H/W Support

- Need to Support:
 - at least two modes:
 - Privileged supervisor mode (aka kernel mode)
 - User mode
 - Different Processor states:
 - Allow for a process to read only; No Write allowed
 - To write, privileged instructions are needed that are only available in supervisor mode
 - Page tables and other state information only accessible in supervisor mode
 - System call exception (e.g., syscall in MIPS):
 - This allows a process to change mode:
 - User mode => supervisor mode (and vice versa)



Memory Protection

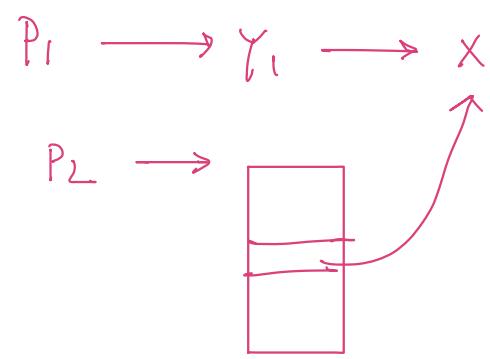
- Each process has its own virtual address space.
- OS can keep the page tables organized so that the independent virtual pages map to disjoint physical pages
 - one process will not be able to access another's data.
- But what if the Process changes the mapping?
- Page tables are placed in the protected address space of the OS
 - So, user process is refrained from changing the page table mapping.
 - But OS is able to modify the page tables.

Memory Sharing and Protection

- OS assists processes for limited information sharing
 - OS can change the Page table as needed
 - The write access bit is used to restrict write
 - Can be changed only by OS

Example:

- P2 wants P1 to access its page
- P2 asks OS to create a page table entry for a virtual page in P1's address space that points to the same physical page that P2 wants to share.
- Any bits that determine the access rights for a page must be included in both the page table and the TLB because the page table is accessed only on a TLB miss.



P₁ → Y₁ → X
P₂ → handle → X

P₁ X কে P₂ এর মাধ্যে share
করতে চায়। P₂ এর page table-এ
X এর জন্য ...
OS handle করবে।

OS middle ground

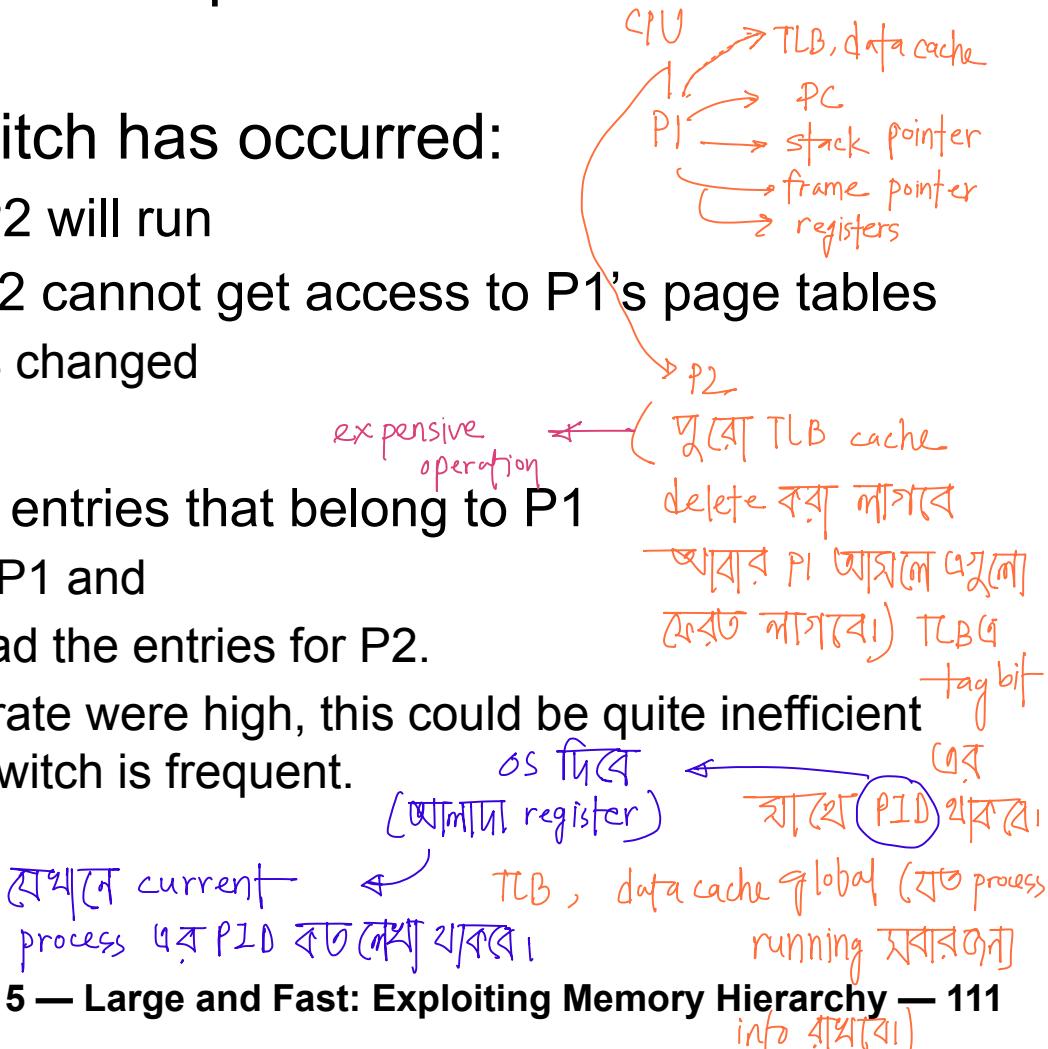
Memory Protection: Context Switch

ম্যাপিং চলাচল শর্ত এবং save
করে রাখা লাগবে।

- Context Switch: A changing of the internal state of the processor to allow a different process to use the processor

- Suppose a context switch has occurred:

- P1 was running; now P2 will run
- OS must ensure that P2 cannot get access to P1's page tables
 - Page Table register is changed
- What about the TLB?
 - OS must clear the TLB entries that belong to P1
 - to protect the data of P1 and
 - to force the TLB to load the entries for P2.
 - If the process switch rate were high, this could be quite inefficient especially if context switch is frequent.



Thread বেশি হলো — context switching বেশি হয়।

data intensive operation বেশি হলো — single thread
চালো perform করবে।

এখন context switching এর ফলয়
CPU কাজ না করে বসে থাকে

memory intensive operation বেশি (database এর মাধ্যে communicate করা
শীত)

— তখন context switching বেশি হলো (problem নেই। কারণ CPU
প্রয়োজন করে না।)

Memory Protection: Context Switch

A common alternative:

extend the virtual address space by adding a *process identifier* or *task identifier*.

The Intrinsity FastMATH has an 8-bit address space ID (ASID) field for this purpose.

- This small field identifies the currently running process
- it is kept in a register loaded by OS when it switches processes.
- The process identifier is concatenated to the tag portion of the TLB
- TLB hit occurs only if both the page number *and* the process identifier match.
- This eliminates the need to clear the TLB in context switch

Similar problems can occur for a cache, since on a process switch the cache will contain data from the running process.

e.g. TLB miss — exception bit on

Exception Enable/Disable

- Suppose we have a page fault exception and OS is handling it
- What will happen, if a second exception occurs?
 - The control unit would overwrite the exception program counter, making it impossible to return to the instruction that caused the page fault!
- We need the ability to **disable** and **enable exceptions**.
 - not ready to receive, যাবি
- When an exception first occurs, the processor sets a bit that disables all other exceptions;
 - exception
যায়তন কুল
ক্ষেত্র যায়বা
সার্ভিস হান্ডলিং
- this could happen at the same time the processor sets the supervisor mode bit.
 - কোম্প পার্যবে না, EPC overwrite (context switching প্রেরণ)
- The OS will then save just enough state to allow it to recover if another exception occurs
 - the exception program counter (EPC) and Cause registers
 - register. লোঞ্চ থাকে হোন line execute করতে
গিয়ে exception হচ্ছে।
 - The operating system can then re-enable exceptions.



Special control registers that help with exceptions, TLB misses, and page faults

Register	CPO register number	Description
EPC	14	Where to restart after exception
Cause	13	Cause of exception
BadVAddr	8	Address that caused exception
Index	0	Location in TLB to be read or written
Random	1	Pseudorandom location in TLB
EntryLo	2	Physical page address and flags
EntryHi	10	Virtual page address
Context	4	Page table address and page number

FIGURE 5.33 MIPS control registers. These are considered to be in coprocessor 0, and hence are read using `mfcc0` and written using `mtcc0`.

When a TLB miss occurs, the MIPS hardware saves the page number of the reference here.

intel processor বিজ্ঞান কাণ্ড করো?

- self study .

TLB Miss in MIPS

- TLB Miss exception invokes OS, which handles the miss in software.
- Control is transferred to 8000 0000hex (TLB Miss Handler Address)
- To find the physical address for the missing page, the TLB miss routine indexes the page table using the page number of the virtual address and the page table register
- To make this indexing fast, MIPS hardware places the address of the Page Table Entry in a special Context Register
- Thus, the first two instructions copy the Context register into the kernel temporary register \$k1 and then load the page table entry from that address into \$k1.
- Recall that \$k0 and \$k1 are reserved for the operating system

TLB Miss in MIPS

TLBmiss:

```
mfc0 $k1,Context      # copy address of PTE into temp $k1
lw    $k1,0($k1)       # put PTE into temp $k1
mtc0 $k1,EntryLo      # put PTE into special register EntryLo
tlbwr                      # put EntryLo into TLB entry at Random
eret                         # return from TLB miss exception
```

- TLB miss handler does not check to see if the page table entry is valid.

If invalid, another and different exception occurs, and OS recognizes the page fault.

- (frequent) TLB miss becomes fast
- at a slight performance penalty for the (infrequent) page fault.

it transfers
control to
8000 0180h

The Memory Hierarchy

The BIG Picture

- Common principles apply at all levels of the memory hierarchy
 - Based on notions of caching
- At each level in the hierarchy
 - Block placement
 - Finding a block
 - Replacement on a miss
 - Write policy

Block Placement

- Determined by associativity
 - Direct mapped (1-way associative)
 - One choice for placement
 - n-way set associative
 - n choices within a set
 - Fully associative
 - Any location
- Higher associativity reduces miss rate
 - Increases complexity, cost, and access time

Finding a Block

Associativity	Location method	Tag comparisons
Direct mapped	Index	1
n-way set associative	Set index, then search entries within the set	n
Fully associative	Search all entries	#entries
	Full lookup table	0

- Hardware caches
 - Reduce comparisons to reduce cost
- Virtual memory
 - Full table lookup makes full associativity feasible
 - Benefit in reduced miss rate

Replacement

- Choice of entry to replace on a miss
 - Least recently used (LRU)
 - Complex and costly hardware for high associativity
 - Random
 - Close to LRU, easier to implement
- Virtual memory
 - LRU approximation with hardware support

Write Policy

- Write-through
 - Update both upper and lower levels
 - Simplifies replacement, but may require write buffer
- Write-back
 - Update upper level only
 - Update lower level when block is replaced
 - Need to keep more state
- Virtual memory
 - Only write-back is feasible, given disk write latency

Sources of Misses

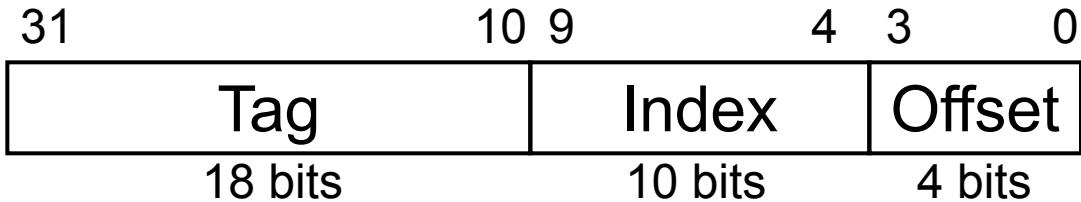
- Compulsory misses (aka cold start misses)
 - First access to a block
- Capacity misses
 - Due to finite cache size
 - A replaced block is later accessed again
- Conflict misses (aka collision misses)
 - In a non-fully associative cache
 - Due to competition for entries in a set
 - Would not occur in a fully associative cache of the same total size

Cache Design Trade-offs

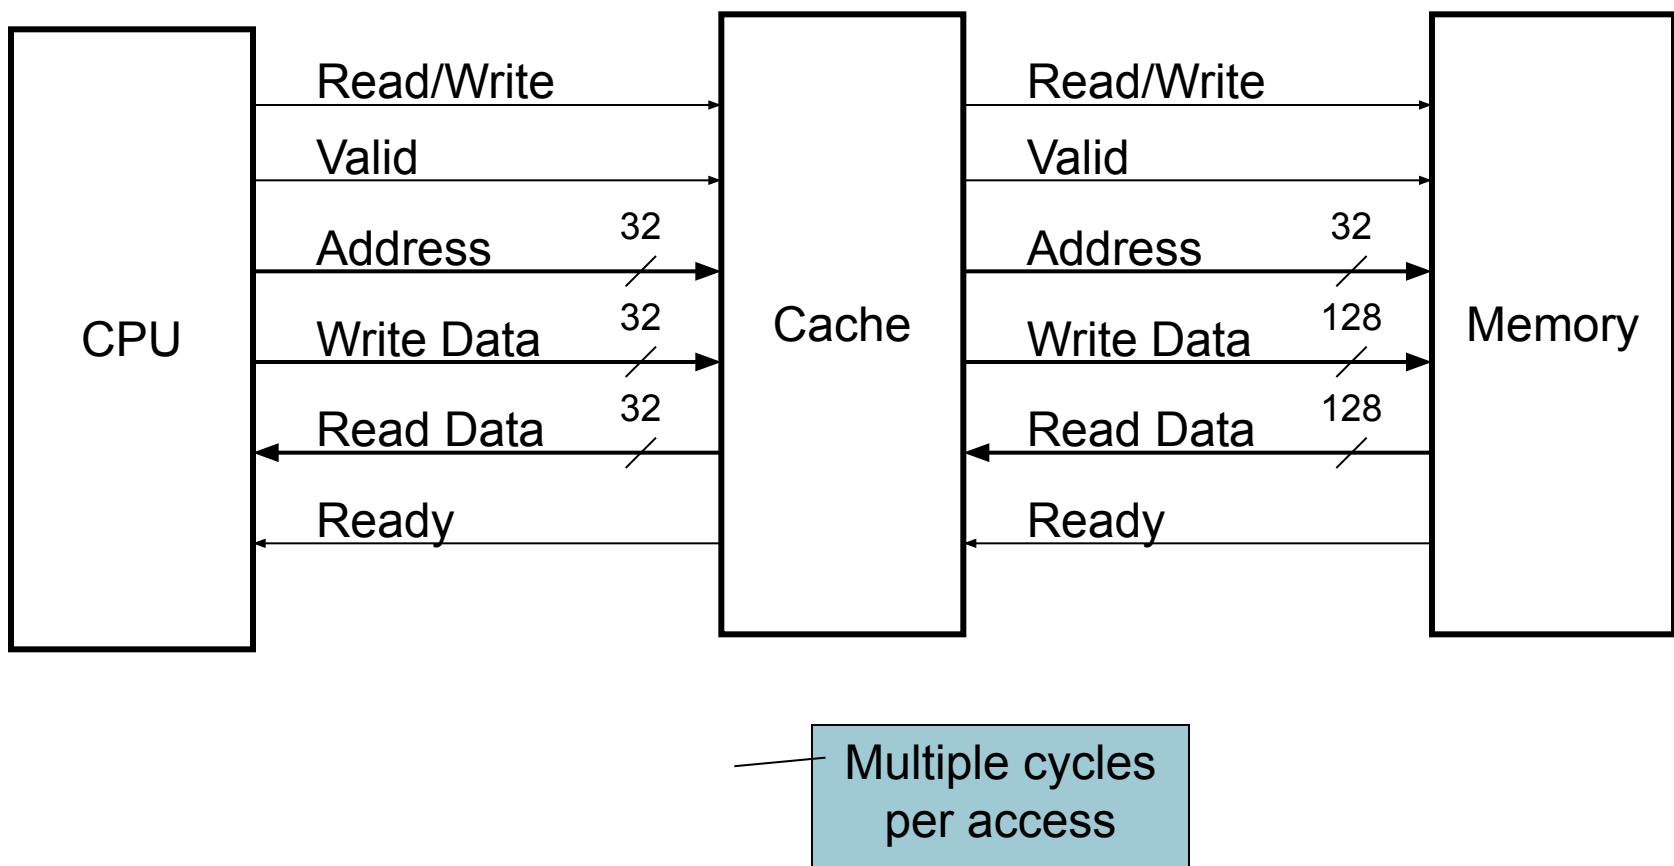
Design change	Effect on miss rate	Negative performance effect
Increase cache size	Decrease capacity misses	May increase access time
Increase associativity	Decrease conflict misses	May increase access time
Increase block size	Decrease compulsory misses	Increases miss penalty. Very large block could increase miss rate.

Cache Control

- Example cache characteristics
 - Direct-mapped, write-back, write allocate
 - Block size: 4 words (16 bytes)
 - Cache size: 16 KB (1024 blocks)
 - 32-bit byte addresses
 - Valid bit and dirty bit per block
 - Blocking cache
 - CPU waits until access is complete

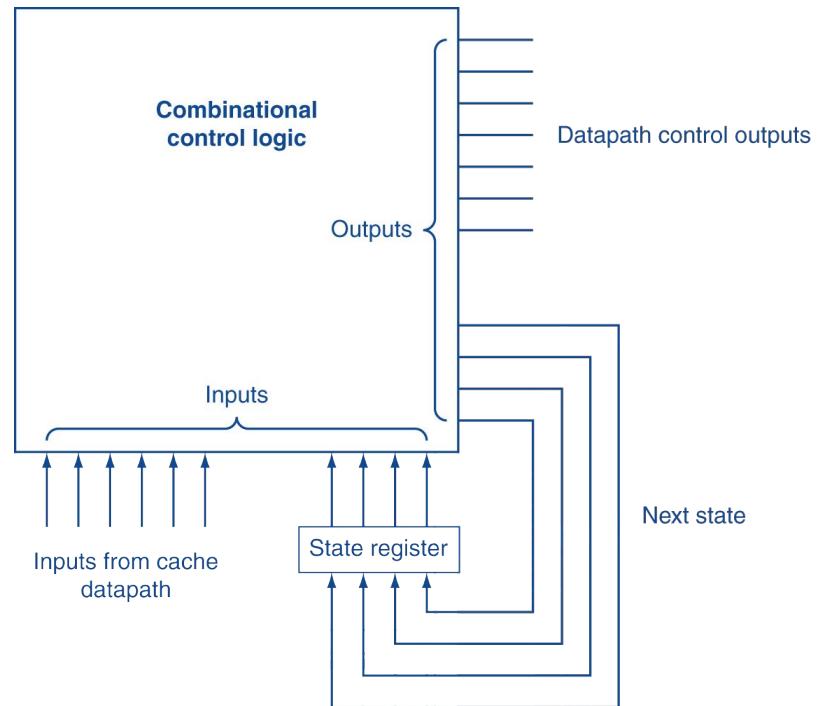


Interface Signals

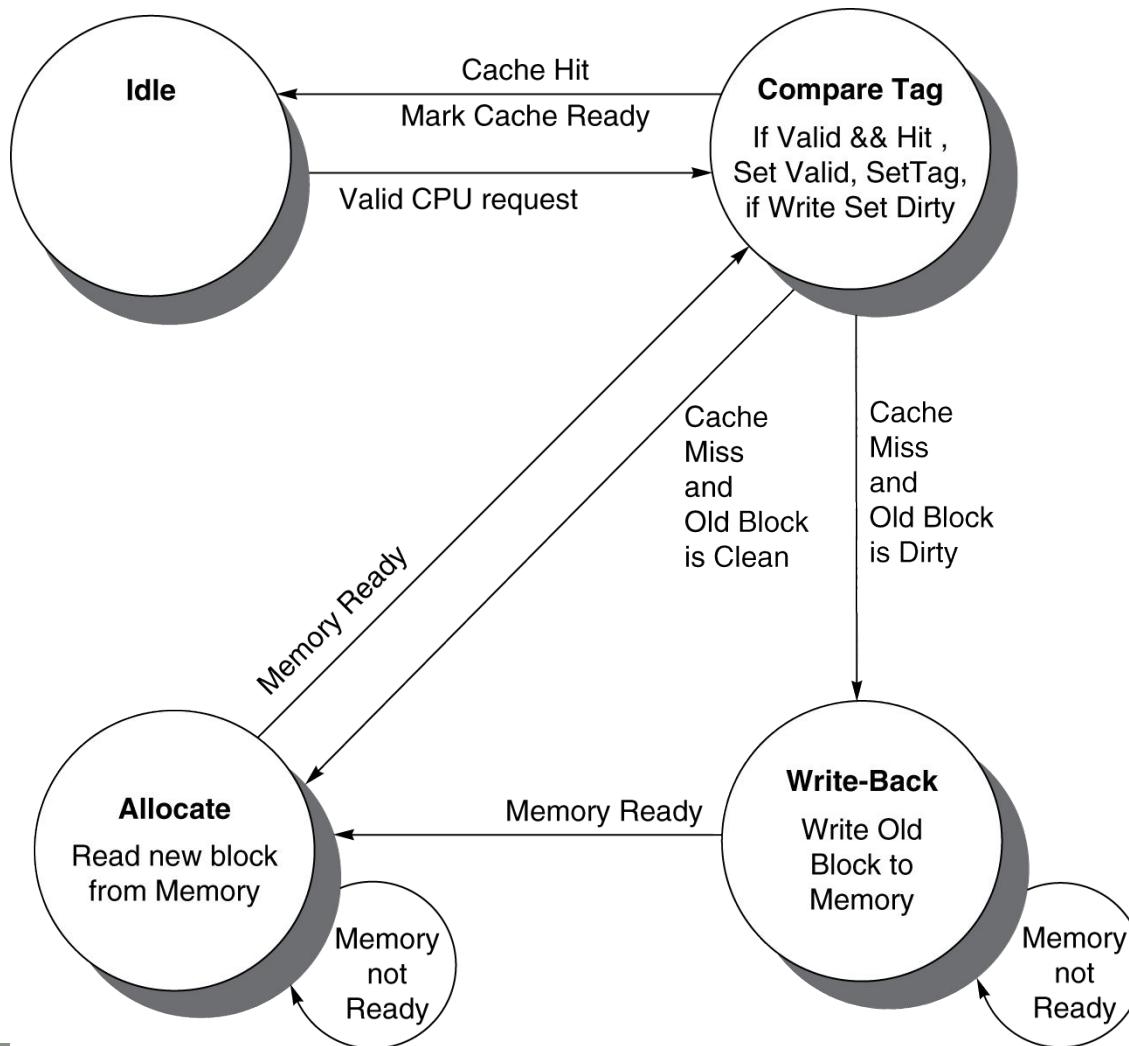


Finite State Machines

- Use an FSM to sequence control steps
- Set of states, transition on each clock edge
 - State values are binary encoded
 - Current state stored in a register
 - Next state = f_n (current state, current inputs)
- Control output signals = f_o (current state)



Cache Controller FSM



Could partition into separate states to reduce clock cycle time